

3fs实验

实验结果

BigDataBench-MPI-Sort对比实验

和其他分布式文件系统对比：

benchmark	3FS运行时间	HDFS运行时间	Lustre系统运行时间
MPI-Sort	33s	40s	5s

使用fio工具进行实验

使用fio工具的原因是3fs本身不使用cache，可以通过fio工具进行直接测试（-direct=1参数）。

读取结果：

```
顺序读：
READ: bw=43.8MiB/s (45.9MB/s), 43.8MiB/s-43.8MiB/s (45.9MB/s-45.9MB/s),
io=7880MiB (8263MB), run=180073-180073msec
随机读：
READ: bw=40.4MiB/s (42.4MB/s), 40.4MiB/s-40.4MiB/s (42.4MB/s-42.4MB/s),
io=7280MiB (7634MB), run=180135-180135msec
```

BigDataBench-MPI-Sort实验

排序时间统计

单机上排序时间：33sec

【五次实验取众数,33sec-36sec之间】

```
OMPI_ALLOW_RUN_AS_ROOT=1 OMPI_ALLOW_RUN_AS_ROOT_CONFIRM=1 mpirun
./mpi_sort /mnt/3fs/wiki-1G /mnt/3fs/sort_data_out
```

实验结果:

```
root@prj2-vm1:~/MPI/MPI_Sort# OMPI_ALLOW_RUN_AS_ROOT=1 OMPI_ALLOW_RUN_AS_ROOT_CONFIRM=1 mpirun ./mpi_sort /mnt/3fs/wiki-1G /mnt/3fs/sort_data_out_2
Thu Jun 12 14:40:49 2025
##folder
process file /mnt/3fs/wiki-1G/lda_wikilw_1
File Size:527079336
File Part Size:514725
read to:58292
1 processes mandates root height of 0
read to:58292
process file /mnt/3fs/wiki-1G/lda_wikilw_2
File Size:171179146
File Part Size:167167
read to:147517
1 processes mandates root height of 0
read to:147517
Total running time:33.000000 sec
Thu Jun 12 14:41:22 2025
root@prj2-vm1:~/MPI/MPI_Sort# OMPI_ALLOW_RUN_AS_ROOT=1 OMPI_ALLOW_RUN_AS_ROOT_CONFIRM=1 mpirun ./mpi_sort /mnt/3fs/wiki-1G /mnt/3fs/sort_data_out
Thu Jun 12 14:02:32 2025
##folder
process file /mnt/3fs/wiki-1G/lda_wikilw_1
File Size:527079336
File Part Size:514725
read to:58292
1 processes mandates root height of 0
read to:58292
process file /mnt/3fs/wiki-1G/lda_wikilw_2
File Size:171179146
File Part Size:167167
read to:147517
1 processes mandates root height of 0
read to:147517
Total running time:33.000000 sec
Thu Jun 12 14:03:05 2025
```

排序结果:

multi_host_file内容:

```
root@prj2-vm1:~/MPI/MPI_Sort# cat multi_host_file
10.66.221.139 slots=2
10.66.221.209 slots=2
```

排序时间:

```
root@prj2-vm1:~/MPI/MPI_Sort# OMPI_ALLOW_RUN_AS_ROOT=1 OMPI_ALLOW_RUN_AS_ROOT_CONFIRM=1 mpirun --hostfile multi_host_file -np 2
./mpi_sort /mnt/3fs/wiki-1G /mnt/3fs/sort_data_out_multi
Thu Jun 12 14:26:33 2025
##folder
process file /mnt/3fs/wiki-1G/lda_wikilw_1
File Size:527079336
File Part Size:514725
Thu Jun 12 14:26:33 2025
##folder
process file /mnt/3fs/wiki-1G/lda_wikilw_1
File Size:527079336
File Part Size:514725
read to:58292
1 processes mandates root height of 0
read to:58292
1 processes mandates root height of 0
read to:58292
read to:58292
process file /mnt/3fs/wiki-1G/lda_wikilw_2
process file /mnt/3fs/wiki-1G/lda_wikilw_2
File Size:171179146
File Size:171179146
File Part Size:167167
File Part Size:167167
read to:147517
read to:147517
1 processes mandates root height of 0
1 processes mandates root height of 0
read to:147517
read to:147517
Total running time:122.000000 sec
Thu Jun 12 14:28:35 2025
Total running time:122.000000 sec
Thu Jun 12 14:28:35 2025
rm: cannot remove 'temp*.txt': No such file or directory
root@prj2-vm1:~/MPI/MPI_Sort#
```

排序结果:

```
head -20 /mnt/3fs/sort_data_out_multi
```

```
root@prj2-vm1:~/MPI/MPI_Sort# head -20 /mnt/3fs/sort_data_out_multi

-2009 -2009 -2009 -2009 -2009 -2009 -2008 -2008 -2008 -2008
-2008 -2008 -2008 -2008 -2008 -2008 -2008 -2008 -2008 -2007
-2007 -2007 -2007 -2007 -2007 -2007 -2007 -2007 -2007 -2007
-2007 -2006 -2006 -2006 -2006 -2006 -2006 -2006 -2006 -2006
-2006 -2006 -2006 -2006 -2006 -2006 -2005 -2005 -2005 -2005
-2005 -2004 -2004 -2004 -2004 -2003 -2003 -2003 -2003 -2002
-2002 -2002 -2002 -24 -24 -24 -1 -1 -1 -1
-1 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0
```

```
tail -20 /mnt/3fs/sort_data_out_multi
```

```
root@prj2-vm1:~/MPI/MPI_Sort# tail -20 /mnt/3fs/sort_data_out_multi
2011 2011 2011 2011 2011 2011 2011 2011 2011 2011
2011 2011 2011 2011 2011 2011 2011 2011 2011 2011
2011 2011 2011 2011 2011 2011 2011 2011 2011 2011
2011 2011 2011 2011 2011 2011 2011 2011 2011 2011
2011 2011 2011 2011 2011 2011 2011 2011 2011 2011
2011 2011 2011 2011 2011 2011 2011 2011 2011 2011
2011 2011 2011 2011 2011 2011 2011 2011 2011 2011
2011 2011 2011 2011 2011 2011 2011 2011 2011 2011
2011 2011 2011 2011 2011 2011 2011 2011 2011 2011
2011 2011 2011 2011 2011 2011 2011 2011 2011 2011
2012 2012 2012 2012 2012 2012 2012 2012 2014 2014 2014
2014 2014 2014 2014 2014 2014 2014 2014 2014 2014
2014 2014 2014 2014 2014 2014 2014 2014 2014 2014
2014 2014 2014 2014 2014 2014 2014 2014 2014 2014
2014 2014 2014 2014 2014 2014 2014 2014 2014 2014
2014 2014 2014 2014 2014 2014 2014 2014 2014 2014
2014 2014 2014 2014 2014 2014 2014 2014 2014 2014
2014 2014 2014 2014 2014 2014 2014 2014 2014 2014
3000 3000 3000 3000 3166 3166 3166 3166 3166 3166
3166 3166 3166 3166 3166 3166 3166 5000 5000 5000
5000 5000 5000 5000 5000 8072 8072 8072 202009 202009
202009 202009 202009 202009 202009 202009 202009 202009 root@prj2-vm1:~/MPI/MPI_Sort#
```

猜想：

受网络影响较大。

关于BigBench-MPI-Sort的配置

BigDataBench的MPI-Sort生成的是文本+数据

```
echo "Preparing MicroBenchmarks data dir"

WORK_DIR="/mnt/3fs"

echo "WORK_DIR=$WORK_DIR data will be put in $WORK_DIR"

cd ../../BigDataGeneratorSuite/Text_datagen/

echo "print data size GB :"
read GB
a=${GB}
L=$((a * 2))
./gen_text_data.sh lda_wiki1w $L 8000 10000 ${WORK_DIR}/wiki-$a"G"
```

lgeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria
lgeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria a
lgeria algeria algeria algeria algeria algeria flight reprint ied algerian algerian algerian algerian algerian algerian algerian
algerian algerian algerian algerian algerian algerian algerian algerian algerian algerian algerian salah korce directors discha
rge status status status author author author author author author author author cash case case case direction libya افعا
icon icon icon icon ghor existed eagle efefef nomads figure crisis tanks football football football football foot نستان
ball football football football football football football football football football regime regime regime ppl election
election election election election election election election election fiction idema centre habibullah habibullah direct autho
link dynasty dynasty dynasty dynasty onepage introduced like like like 300px 300px land_km2 land_km2 archivedate archivedate ar
chivedate archivedate teams teams teams régiment catégorie buildings sweden subsequently subsequently belgian withdraw stripe ce
ntury century century century century century century century century century century politician politic
ian politician point point romanians adrar adrar adrar biography biography hour africa africa africa africa africa africa afic
a africa africa africa depth consisted extended wake sheberghan august august august august august august august august
ugust across across seal seal seal seal seal seal seal seat declared declared declared declared london london gen
geo geo geo geo get interior interior squad kevin pashtun seeing freename freename abroad bot boy targets amanullah
trends down down down represented represented motto motto motto motto mott troop dostum personal personal personal titula
r pushpin pushpin pushpin pushpin pushpin pushpin pushpin pushpin pushpin pushpin pushpin pushpin pushpin pushpin jpg jp
g jpg jpg jpg jpg jpg jpg jpg jpg berlin spokesman highest highest gain additional http http http http http http h
ttp ht
http ht
tp speaker controlled controlled key key key indo indo nytimes battles battles battles battles descriptions descriptions descrip
tions biggest region region region region region region region region region region special special special including inc
luding including including including including links links links links links links links links links withdrew withdrew
similar empty mausoleum canadians minorities league league league promoted american american american american america
n american american heavy evening helmand helmand helmand goals4 sheikh timeline religious religious religious religious religio
us bulk drama death death death death death death death death death develop develop develop metric metri
c colours iaaf romania transition taliban taliban taliban taliban taliban taliban taliban taliban taliban taliban taliba
n taliban taliban taliban taliban taliban taliban taliban still still still ended armies trans turns sovereign night cas
bah casbah gns gns gns cooperation cooperation peshawar insurgents insurgents insurgents insurgents return return ano
ther another another another another another another disputes disputed tauris birth birth birth birth birth birth
birth birth birth birth birth birth production production production 1998 1998 1998 1998 1998 1998 1999 1999 1999 1999 1994 1994
1994 1994 1994 1995 1995 1995 1995 1995 1992 1992 1992 1992 1992 1993 1993 1990 1990 1991 1991 1991 championships championships
championships sportspeople worth worth details details details problem xinhuanet zimbabwe off off greco genetic abbr abbr abbr
entities enforcement could could books books books books books books books books books books books books books books
oruç ibid flowing victor video video video terrorists assisted lady currentclub sangin organisations skanderbeg huge company co
mpany company company company more more more more more more more more more more more more more more more abuses being
being being being being turned quarter quarter quarter following following following following following following following fol
lowing following geography geography abdullah abdullah connecting chair chain mirror crownextent nature nature iata master need

BigDataBench数据生成结果:

d end end end end my my mr mr mr ms mm mm mm mm mc mc mf me leader1 leaders leaders watan 80 minutes minutes 1981 1981 1983 1983
1982 1982 1982 1982 1985 1984 1987 1986 annaba works center center center center center center center conquered c
onquered areas areas areas kamboj kia afp afp sword sword sword forward forward britannica per per per pew precipitation precipi
tation precipitation precipitation y moment ireland create supply supply supply koh blood bulgarian bulgarian saudi fb fb fb fb
fb fb fb same same same middle middle destinations je killing investigations americans americans finance finance balkh balk
h 1996 1996 1996 1996 1996 id id id id id id id ie io il im iv stop plurality plurality isa iso decree decree never de
lvniñ dmy dmy inner expansion historia addition wardak playername she she she ece missile effort effort requires required requir
ed required run run run march march march march march march march march march march march march march march march march
bgovernor governor governor governor governor governor described shifted wildlife 1870 1873 asp foot witnessed real kenya kenya
gender opened es es es es es es et et et et et ei en eo el el el el el birds goal goal structure structure december
december december december december december december kemp change change change appears appears explosion supporting g
lobal minor eulma graduated insurgent arab arab arab arab valign sē adding blamed norwegian confirmed vlachs 100px 100px 100px r
rethi rrethi hosted indiana secure established established established established established established establish
ed established alliance alliance alliance takhar documents documents jesu iht agriculture agriculture frequency began
began road road road road ministers mehmed mehmet internationally index index index index index gardez garden summit points poin
ts codes codes codes codes codes ambush battalion battalion battalion battalionlocation location location location loc
ation location levels sea sea sea see see see see see see see see see see see set set set au ar ar ar ar ar
ar ar ar ap al al al al al al al al al al am ah af af ad ad notoc names elwatan duty costs djanel influence in
fluence shahi shahr oxus oxus 2 bundesliga revised records match
es decades kabul kabul kabul kabul kabul kabul kabul kabul kabul kabul kabul believed believes studies studies curren
nt current current type type type type type type type type type type type type type type type type type type
ariana ariana life best best besa thessaloniki gauge engagement machine artillery artillery regional regional yellow car
rier grounds why who who who who who who who who who who who who who who who who 1500 1500 unknown footnotes footnot
es footnotes footnotes footnotes footnotes footnotes footnotes footnotes footnotes footnotes footnotes footnotes footn
otes footnotes footnotes footnotes footnotes footnotes footnotes footnotes ambassadors kunduz chose chose nick rally prominent a
rghandab higher higher algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria a
lgeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria
lgeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria a
lgeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria algeria a
lgeria algeria algeria algeria algeria algeria flight reprint ied algerian algerian algerian algerian algerian algerian algerian
algerian algerian algerian algerian algerian algerian algerian algerian algerian algerian algerian algerian algerian algerian
rge status status status status author author author author author author author author cash case case case direction libya اف
icon icon icon icon ghor existed eagle efefef nomads figure crisis tanks football football football football football foot تسان
ball football football football football football football football football football regime regime regime ppl election
election election election election election election election fiction idema centre habibullah habibullah direct autho
rink dynasty dynasty dynasty dynasty onepage introduced like like like 300px 300px land km2 land km2 archivedate archivedate ar
chivedate archivedate teams teams teams teams régiment catégorie buildings sweden subsequently subsequently belgian withdraw stripe ce

其他问题

排序文本预期生成的是1个G，但是为什么输出的是666M？

```
root@prj2-vm1:~/MPI/MPI_Sort# sudo sh genData_sort.sh
Preparing MicroBenchmarks data dir
WORK_DIR=/mnt/3fs data will be put in /mnt/3fs
print data size GB :
1
Thu Jun 12 06:12:52 UTC 2025

Thu Jun 12 12:12:35 UTC 2025
root@prj2-vm1:/mnt/3fs# du -sh wiki-1G/
666M      wiki-1G/
```

fio实验

顺序读：

```
fio -directory=/mnt/3fs/ -numjobs=1 -fallocate=none -iodepth=2 -
ioengine=libaio -direct=1 -rw=read -bs=4M --group_reporting -size=100M -
time_based -runtime=180 -name=2depth_128file_4M_direct_read_bw
```

```
root@prj2-vm1:~/MPI/MPI_Sort# fio -directory=/mnt/3fs/ -numjobs=1 -fallocate=none -iodepth=2 -ioengine=libaio -direct=1 -rw=read -bs=4M --group_reporting -size=100M -time_based -runtime=180 -name=2depth_128file_4M_direct_read_bw
2depth_128file_4M_direct_read_bw: (groupid=0, jobs=1): err= 0: pid=7636: Thu Jun 12 15:08:15 2025
  read: IOPS=10, BW=43.8MiB/s (45.9MB/s)(7880MiB/180073msec)
    slat (usec): min=54, max=7522, avg=149.70, stdev=313.45
    clat (msec): min=47, max=1900, avg=182.63, stdev=162.89
    lat (msec): min=48, max=1900, avg=182.78, stdev=162.89
  clat percentiles (msec):
    | 1.00th=[ 80], 5.00th=[ 90], 10.00th=[ 97], 20.00th=[ 109],
    | 30.00th=[ 117], 40.00th=[ 127], 50.00th=[ 138], 60.00th=[ 148],
    | 70.00th=[ 165], 80.00th=[ 199], 90.00th=[ 296], 95.00th=[ 435],
    | 99.00th=[ 1003], 99.50th=[ 1301], 99.90th=[ 1687], 99.95th=[ 1905],
    | 99.99th=[ 1905]
  bw ( KiB/s): min= 4724, max=81920, per=100.00%, avg=46489.33, stdev=20017.26, samples=345
  iops        : min=   1, max=  20, avg=10.94, stdev= 4.88, samples=345
  lat (msec)  : 50=0.05%, 100=12.39%, 250=74.42%, 500=9.24%, 750=2.08%
  lat (msec)  : 1000=0.86%, 2000=0.96%
  cpu         : usr=0.01%, sys=0.13%, ctx=2309, majf=0, minf=2059
  IO depths   : 1=0.1%, 2=99.9%, 4=0.0%, 8=0.0%, 16=0.0%, 32=0.0%, >=64=0.0%
  submit      : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%, 64=0.0%, >=64=0.0%
  complete    : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%, 64=0.0%, >=64=0.0%
  issued rwts: total=1970,0,0,0 short=0,0,0,0 dropped=0,0,0,0
  latency     : target=0, window=0, percentile=100.00%, depth=2

Run status group 0 (all jobs):
  READ: bw=43.8MiB/s (45.9MB/s), 43.8MiB/s-43.8MiB/s (45.9MB/s-45.9MB/s), io=7880MiB (8263MB), run=180073-180073msec
```

随机读：

```
fio -directory=/mnt/3fs/ -numjobs=1 -fallocate=none -iodepth=2 -
ioengine=libaio -direct=1 -rw=randread -bs=4M --group_reporting -
size=100M -time_based -runtime=180 -name=2depth_128file_4M_direct_read_bw
```

```

root@prj2-vm1:~/MPI/MPI_Sort# fio -directory=/mnt/3fs/ -numjobs=1 -fallocate=none -iodepth=2 -ioengine=libaio -direct=1 -rw=randread -bs=4M --group_reporting -size=100M -time_based -runtime=180 -name=2depth_128file_4M_direct_read_bw
2depth_128file_4M_direct_read_bw: (groupid=0, jobs=1): err= 0: pid=7644: Thu Jun 12 15:17:41 2025
  read: IOPS=10, BW=40.4MiB/s (42.4MB/s) (7280MiB/180135msec)
    slat (usec): min=66, max=10933, avg=176.56, stdev=440.81
    clat (msec): min=68, max=1940, avg=197.71, stdev=177.94
      lat (msec): min=68, max=1940, avg=197.89, stdev=177.94
    clat percentiles (msec):
      | 1.00th=[ 91], 5.00th=[ 101], 10.00th=[ 107], 20.00th=[ 117],
      | 30.00th=[ 127], 40.00th=[ 136], 50.00th=[ 146], 60.00th=[ 161],
      | 70.00th=[ 178], 80.00th=[ 213], 90.00th=[ 330], 95.00th=[ 464],
      | 99.00th=[ 986], 99.50th=[ 1418], 99.90th=[ 1938], 99.95th=[ 1938],
      | 99.99th=[ 1938]
  bw ( KiB/s): min= 3065, max=81920, per=100.00%, avg=43143.19, stdev=18905.90, samples=343
  iops       : min=   0, max=  20, avg=10.15, stdev= 4.62, samples=343
  lat (msec) : 100=5.22%, 250=79.51%, 500=11.04%, 750=2.36%, 1000=0.93%
  lat (msec) : 2000=0.93%
  cpu        : usr=0.02%, sys=0.13%, ctx=2165, majf=0, minf=2058
  IO depths  : 1=0.1%, 2=99.9%, 4=0.0%, 8=0.0%, 16=0.0%, 32=0.0%, >=64=0.0%
  submit     : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%, 64=0.0%, >=64=0.0%
  complete   : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%, 64=0.0%, >=64=0.0%
  issued rwts: total=1820,0,0,0 short=0,0,0,0 dropped=0,0,0,0
  latency    : target=0, window=0, percentile=100.00%, depth=2

Run status group 0 (all jobs):
  READ: bw=40.4MiB/s (42.4MB/s), 40.4MiB/s-40.4MiB/s (42.4MB/s-42.4MB/s), io=7280MiB (7634MB), run=180135-180135msec

```

具体部署事项

写在前面

没有真实RDMA设备采用softRoCE方式时，外部必须是**Generic**内核的完整**Linux**环境，不管是编译部署还是docker部署，否则没有rxr模块。

非常容易忽略的点，如果是硬件指令集没有avx512的时候，建议直接换有avx512机器。根据实践，使用虚拟机上是无法正常运行的（无论使用哪一种部署方式），哪怕给4G内存，128G的磁盘。在本地虚拟机上硬部署的结果（这个也有可能是本地的硬件带不动，但是由于之前服务器上的lxc容器是没有内核的，所以所有的尝试都在本地，大失败）：

1. 按照官方文件编译配置，编译成功，但是无法正常使用，尤其是Storage节点无法启动，直接卡死挂掉。

```

root@meta-01:/home/wang# /opt/3fs/bin/admin_cli -cfg /opt/3fs/etc/admin_cli.toml --config.mgmt_client.mgmt
td_server_addresses '["RDMA://192.168.174.130:8000"]' "list-nodes"
bash: line 1: /usr/sbin/ibdev2netdev: No such file or directory

```

Id	Type	Status	Hostname	Pid	Tags	LastHeartbeatTime	ConfigVersion	ReleaseVersion
1	MGMTD	PRIMARY_MGMTD	meta-01	2603	[]	N/A	2(UPTODATE)	250228-dev-1-99
9999-91bf3cf36								
100	META	HEARTBEAT_CONNECTED	meta-01	3456	[]	2025-06-08 23:29:44	1(UPTODATE)	250228-dev-1-99
9999-91bf3cf36								

2. 按照docker方式部署，部署成功，挂载成功，但是storage异常退出，退出码351，为指令不兼容。使用-avx2版本的结果：通官方编译结果，直接卡死，挂掉，而且难以收集异常退出信息。

```

root@node1:/home/wang/Desktop/m3fs/m3fs# mount | grep 3fs
hf3fs.open3fs on /mnt/3fs type fuse.hf3fs (rw,nosuid,nodev,relatime,

```

```

root@node1:/home/wang/Desktop/m3fs/m3fs# docker ps -a

```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
36735725b036	open3fs/3fs:20250410	"/opt/3fs/bin/hf3fs_..."	3 minutes ago	Up 2 minutes		3fs-client
231e5a9c310b	open3fs/3fs:20250410	"/opt/3fs/bin/storag..."	3 minutes ago	Exited (132) 3 minutes ago		3fs-storage
9bcd8f84f659	open3fs/3fs:20250410	"/opt/3fs/bin/meta_m..."	3 minutes ago	Up 3 minutes		3fs-meta
f7e796fb322	open3fs/3fs:20250410	"/opt/3fs/bin/mgmt_d..."	3 minutes ago	Up 3 minutes		3fs-mgmt
191a45d3c449	open3fs/3fs:20250410	"/opt/3fs/bin/monito..."	3 minutes ago	Up 3 minutes		3fs-monitor
d774a4bb05b8	open3fs/grafana:12.0.0	"/run.sh"	3 minutes ago	Up 3 minutes		3fs-grafana
1f7ef8ee39f8	open3fs/clickhouse:25.1-jammy	"/entrypoint.sh"	3 minutes ago	Up 3 minutes		3fs-clickhouse
b5649b33c49f	open3fs/foundationdb:7.3.63	"/usr/bin/tini -g --..."	3 minutes ago	Up 3 minutes		3fs-fdb

完全没成功的lscpu，只有avx2:

```

root@meta:/home/wang/Desktop# lscpu | grep avx512
root@meta:/home/wang/Desktop# lscpu | grep avx2
Flags: fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush mmx fxsr
sse sse2 ss syscall nx pdpe1gb rdtscp lm constant_tsc arch_perfmon rep_good nopl xtopology tsc_reliable nonstop_tsc cpuid tsc_
nwnw_freq pni pclmulqdq ssse3 fma cx16 sse4_1 sse4_2 x2apic movbe popcnt aes xsave avx f16c rdrand hypervisor lahf_lm abm 3dnow
prefetch pti ssbd ibrs ibpb stibp fsgsbase tsc_adjust bmi1 avx2 smep bmi2 erms invpcid rdseed adx snap clflushopt clwb sha_ni
saveopt xsavec xgetbv1 xsaves avx_vnni arat umip gfni vaes vpclmulqdq rdpid movdiri movdir64b fsrm md_clear serialize flush_l1d
arch_capabilities

```


成功的lscpu, 有avx512:

```
dbms-stu-02@teacher-PowerEdge-M640:~$ lscpu | grep avx512
Flags: fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush dts acpi mmx fxsr sse sse2 ss ht tm pbe syscall nx pdpe1gb rdtscp lm constant_tsc art arch_perfmon pebs bts rep_good nopl xtopology nonstop_tsc cpuid aperfmperf pni pclmulqdq dtes64 monitor ds_cpl vmx smx est tm2 ssse3 sdbg fma cx16 xtpr pdcm pcid dca sse4_1 sse4_2 x2apic movbe po
pcnt tsc_deadline_timer aes xsave avx f16c rdrand lahf_lm abm 3dnowprefetch cpuid_fault epb cat_l3 cdp_l3 invpcid_single pti intel_ppin ssbd mba ibrs ibpb stibp tpr_shadow vnmi flexpriority ept vpid fsgsbase tsc_adjust bmi1 hle avx2 smep bmi2 erms invpcid rtm cqm mpx rdt_a avx512f avx512dq rdseed adx smap clflushopt clwb intel_pt avx512cd avx512bw avx512vl xsaveopt xsavec xgetbv1 xsavec
saves_cqm llc cqm_occup_llc cqm_mbm_total cqm_mbm_local dtherm ida arat pln pts pku ospke md_clear flush_l1d
```

部署方式一：编译文件

按照官方文件进行，编译安装。问题是编译安装时间过长，而且配置文件的修改需要自己手动修改，有些冗余。好处是知道各种结构，搭配配置文件可塑性强。

编译结果：

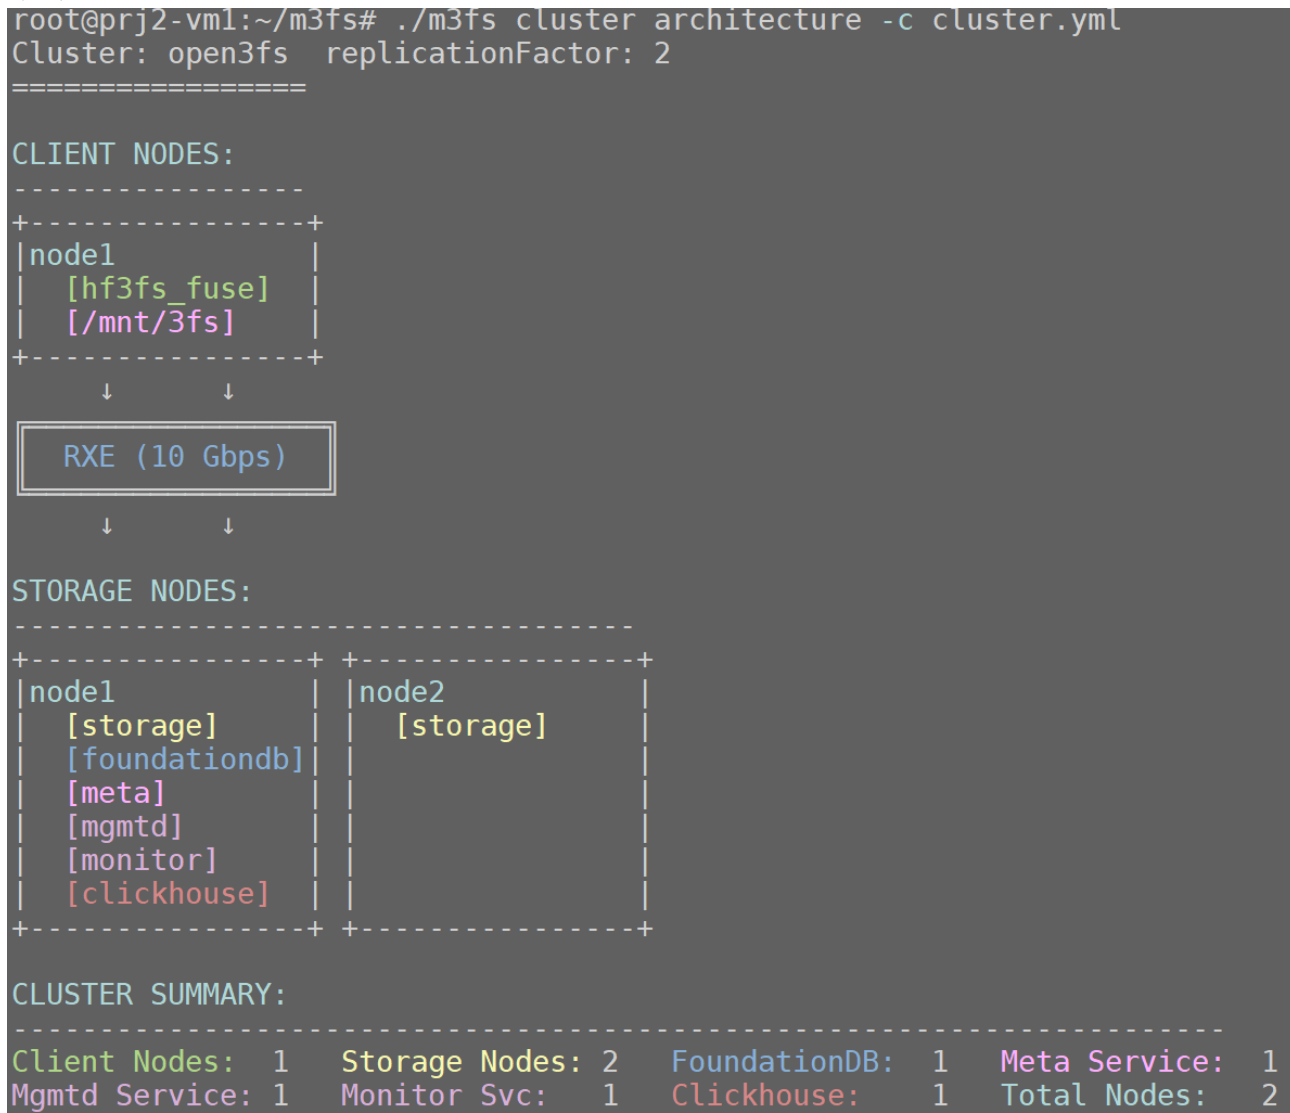
```
[ 97%] Built target test_storage_service
[ 97%] Building CXX object tests/storage/store/CMakeFiles/test_storage_store.dir/...
[ 98%] Building CXX object tests/storage/store/CMakeFiles/test_storage_store.dir/...
[ 98%] Building CXX object tests/storage/store/CMakeFiles/test_storage_store.dir/...
[ 98%] Building CXX object tests/storage/store/CMakeFiles/test_storage_store.dir/...
[ 98%] Building CXX object tests/storage/store/CMakeFiles/test_storage_store.dir/...
[ 98%] Building CXX object tests/storage/store/CMakeFiles/test_storage_store.dir/...
[ 98%] Linking CXX executable ../../test_storage_store
[ 98%] Built target test_storage_store
[ 98%] Building CXX object tests/storage/sync/CMakeFiles/test_storage_sync.dir/...
[ 98%] Building CXX object tests/storage/sync/CMakeFiles/test_storage_sync.dir/...
[ 98%] Linking CXX executable ../../test_storage_sync
[ 98%] Built target test_storage_sync
[ 98%] Built target test_mgmt
[100%] Building CXX object tests/migration/CMakeFiles/test_migration.dir/TestM...
[100%] Linking CXX executable ../test_migration
[100%] Built target test_migration
[100%] Built target follybenchmark
[100%] Building CXX object benchmarks/storage_bench/CMakeFiles/storage_bench.dir/...
[100%] Linking CXX executable ../../bin/storage_bench
[100%] Built target storage_bench
[root@bec40f7d0865 3fs]#
```

可执行文件：

```
[root@bec40f7d0865 bin]# ls
admin_cli      hf3fs_fuse_main  mgmtd_main      monitor_collector_main  storage_bench
hf3fs-admin   meta_main        migration_main   simple_example_main     storage_main
[root@bec40f7d0865 bin]# pwd
/root/3fs/build/bin
```

部署方式二：docker部署

集群结构：



来自：<https://github.com/open3fs/m3fs>

这里的主要问题就是在lxc的Ubuntu容器里对网络的要求有些奇怪：foundationDB的docker中需要ipv4，否则挂掉（无法解析）；同时clickhouse的docker中又要求不能完全直接关闭ipv6。但是在直接的Ubuntu虚拟机中没有这些奇怪的毛病。

其他的就是保证端口没有被占用即可。

```
root@prj2-vm1:~/MPI/MPI_Sort# ip -6 addr
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 state UNKNOWN qlen 1000
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
root@prj2-vm1:~/MPI/MPI_Sort# ip addr
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: enp5s0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP group default qlen 1000
    link/ether 00:16:3e:13:2a:4d brd ff:ff:ff:ff:ff:ff
    inet 10.66.221.139/24 metric 100 brd 10.66.221.255 scope global dynamic enp5s0
        valid_lft 3135sec preferred_lft 3135sec
3: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN group default
    link/ether 02:42:c4:0f:5c:6c brd ff:ff:ff:ff:ff:ff
    inet 172.17.0.1/16 brd 172.17.255.255 scope global docker0
        valid_lft forever preferred_lft forever
```

部署成功界面

```
root@prj2-vm1:~/MPI/MPI_Sort# docker ps -a
CONTAINER ID   IMAGE                                COMMAND                                CREATED        STATUS        PORTS        NAMES
dbcb34364976   open3fs/3fs:20250410               "/opt/3fs/bin/hf3fs_..."          9 hours ago   Up 9 hours   3fs-client   3fs-client
d0c274413e94   open3fs/3fs:20250410               "/opt/3fs/bin/storag..."          9 hours ago   Up 9 hours   3fs-storage   3fs-storage
6ea22dfc8b6c   open3fs/3fs:20250410               "/opt/3fs/bin/meta_m..."          9 hours ago   Up 9 hours   3fs-meta      3fs-meta
26ec5dc83e3d   open3fs/3fs:20250410               "/opt/3fs/bin/mgmt_d..."          9 hours ago   Up 9 hours   3fs-mgmt      3fs-mgmt
45c1f9e127dd   open3fs/3fs:20250410               "/opt/3fs/bin/monito..."          9 hours ago   Up 9 hours   3fs-monitor   3fs-monitor
1ac5c970fe0b   open3fs/grafana:12.0.0              "/run.sh"                            9 hours ago   Up 9 hours   3fs-grafana   3fs-grafana
502fe62d9abb   open3fs/clickhouse:25.1-jammy       "/entrypoint.sh"                    9 hours ago   Up 9 hours   3fs-clickhouse 3fs-clickhouse
6e1d0e9f6021   open3fs/foundationdb:7.3.63         "/usr/bin/tini -g --..."          9 hours ago   Up 26 minutes 3fs-fdb        3fs-fdb
```

另一个storage节点的部署:

```
root@prj2-vm2:~# docker ps -a
CONTAINER ID   IMAGE                                COMMAND                                CREATED        STATUS        PORTS        NAMES
5ea6383ca188   open3fs/3fs:20250410               "/opt/3fs/bin/storag..."          8 hours ago   Up 8 hours   3fs-storage   3fs-storage
```

挂载目录

```
mount | grep 3fs
```

读写方式就和直接读写本地文件系统一样【fuse的方式，便捷，性能会差点：https://github.com/deepseek-ai/3FS/blob/main/docs/design_notes.md#limitations-of-fuse】。当然还有别的方式：https://github.com/deepseek-ai/3FS/blob/main/docs/design_notes.md#asynchronous-zero-copy-api

```
root@prj2-vm1:~/m3fs# mount | grep 3fs
hf3fs.open3fs on /mnt/3fs type fuse.hf3fs (rw,nosuid,nodev,relatime,user_id=0,group_id=0,default_permissions,allow_other,max_rea
d=1048576)
```

挂载大小以及FileSystem:

```
root@prj2-vm1:~/m3fs# df /mnt/3fs/ -h
Filesystem      Size  Used Avail Use% Mounted on
hf3fs.open3fs   181G   88G   93G   49% /mnt/3fs
root@prj2-vm1:~/m3fs# df / -h
Filesystem      Size  Used Avail Use% Mounted on
/dev/root       91G   59G   32G   66% /
```