

中国科学院大学计算机学院专业选修课（硕博通用课程）

# GPU架构与编程

## 第五课：CUDA编程（四）

赵地

中科院计算所

2024年秋季学期

## 讲授内容

- **Related Programming Models: OpenCL**
- **Related Programming Models: OpenACC**
- **Multi-GPU: OpenMP**
- **Related Programming Models: MPI**

# 讲授内容： Related Programming Models: OpenCL

## ① OpenCL Data Parallelism Model

## ② OpenCL Device Architecture

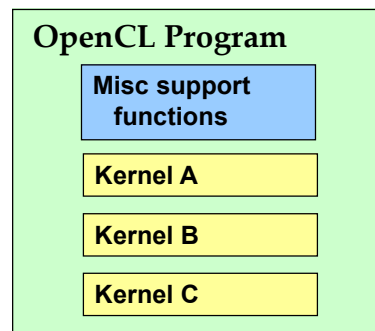
## ③ OpenCL Host Code

## Background

- OpenCL was initiated by Apple and maintained by the Khronos Group (also home of OpenGL) as an industry standard API
  - For cross-platform parallel programming in CPUs, GPUs, DSPs, FPGAs,...
- OpenCL draws heavily on CUDA
  - Easy to learn for CUDA programmers
- OpenCL host code is much more complex and tedious due to desire to maximize portability and to minimize burden on vendors

## OpenCL Programs


- An OpenCL “program” is a C program that contains one or more “kernels” and any supporting routines that run on a target device
- An OpenCL kernel is the basic unit of parallel code that can be executed on a target device



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## OpenCL Execution Model

- Integrated host+device app C program
  - Serial or modestly parallel parts in host C code
  - Highly parallel parts in device SPMD kernel C code



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Mapping between OpenCL and CUDA data parallelism model concepts

OpenCL Parallelism Concept	CUDA Equivalent
host	host
device	device
kernel	kernel
host program	host program
NDRange (index space)	grid
work item	thread
work group	block

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## OpenCL Kernels


- Code that executes on target devices
- Kernel body is instantiated once for each work item
  - An OpenCL work item is equivalent to a CUDA thread
- Each OpenCL work item gets a unique index

```
__kernel void vadd(__global const float *a,
                  __global const float *b,
                  __global float *result)
{
    int id = get_global_id(0);
    result[id] = a[id] + b[id];
}
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Array of Work Items

- An OpenCL kernel is executed by an array of work items
- All work items run the same code (SPMD)
- Each work item can call `get_global_id()` to get its index for computing memory addresses and make control decisions



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Work Groups: Scalable Cooperation

- Divide monolithic work item array into work groups
  - Work items within a work group cooperate via **shared memory and barrier synchronization**
  - Work items in different work groups cannot cooperate
- OpenCL counter part of CUDA Thread Blocks

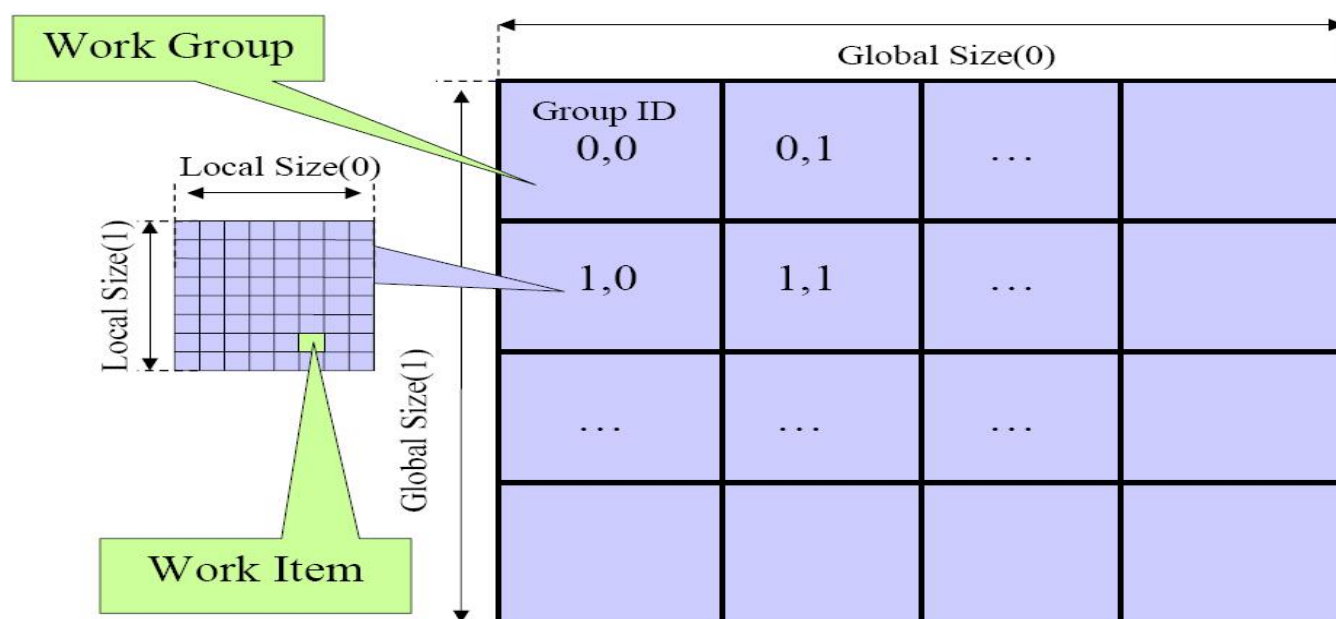
The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## OpenCL Dimensions and Indices

OpenCL API Call	Explanation	CUDA Equivalent
<code>get_global_id(0);</code>	global index of the work item in the x dimension	<code>blockIdx.x*blockDim.x+threadIdx.x</code>
<code>get_local_id(0)</code>	local index of the work item within the work group in the x dimension	<code>threadIdx.x</code>
<code>get_global_size(0);</code>	size of NDRange in the x dimension	<code>gridDim.x*blockDim.x</code>
<code>get_local_size(0);</code>	Size of each work group in the x dimension	<code>blockDim.x</code>

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Multidimensional Work Indexing



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## OpenCL Data Parallel Model Summary

- Parallel work is submitted to devices by launching kernels
- Kernels run over global dimension index ranges (NDRange), broken up into “work groups”, and “work items”
- Work items executing within the same work group can synchronize with each other with barriers or memory fences
- Work items in different work groups can’t sync with each other, except by terminating the kernel


The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## 讲授内容： Related Programming Models: OpenCL

- ① OpenCL Data Parallelism Model
- ② OpenCL Device Architecture
- ③ OpenCL Host Code

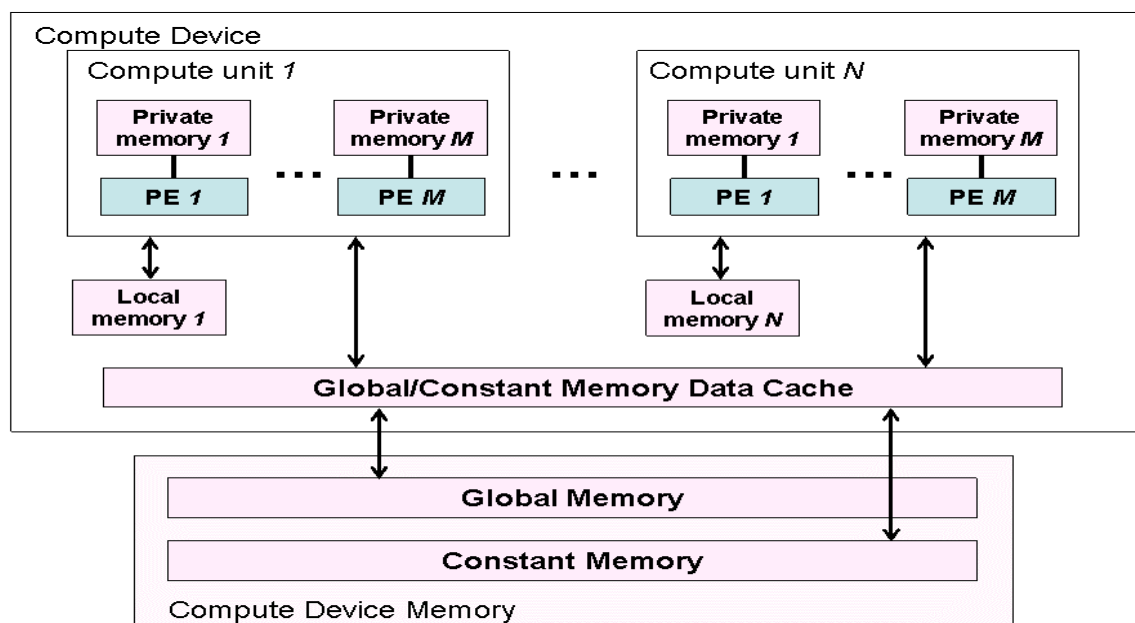
## OpenCL Hardware Abstraction

- OpenCL exposes CPUs, GPUs, and other Accelerators as “devices”
- Each device contains one or more “compute units”, i.e. cores, Streaming Multicprocessors, etc...
- Each compute unit contains one or more SIMD “processing elements”, (i.e. SP in CUDA)



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## OpenCL Device Architecture



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.




## OpenCL Device Memory Types

Memory Type	Host access	Device access	CUDA Equivalent
global memory	Dynamic allocation; Read/write access	No allocation; Read/write access by all work items in all work groups, large and slow but may be cached in some devices.	global memory
constant memory	Dynamic allocation; read/write access	Static allocation; read- only access by all work items.	constant memory
local memory	Dynamic allocation; no access	Static allocation; shared read-write access by all work items in a work group.	shared memory
private memory	No allocation; no access	Static allocation; Read/write access by a single work item.	registers and local memory

## OpenCL Context

- Contains one or more devices
- OpenCL device memory objects are associated with a context, not a specific device



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

# 讲授内容： Related Programming Models: OpenCL

- ① OpenCL Data Parallelism Model
- ② OpenCL Device Architecture
- ③ OpenCL Host Code

## OpenCL Context

- Contains one or more devices
- OpenCL memory objects are associated with a context, not a specific device
- `clCreateBuffer()` is the main data object allocation function
  - error if an allocation is too large for any device in the context
- Each device needs its own work queue(s)
- Memory copy transfers are associated with a command queue (thus a specific device)

## OpenCL Context Setup Code (simple)

```
cl_int clerr = CL_SUCCESS;

cl_context clctx = clCreateContextFromType(0,
CL_DEVICE_TYPE_ALL, NULL, NULL, &clerr);

size_t parmsz;

clerr = clGetContextInfo(clctx, CL_CONTEXT_DEVICES, 0, NULL,
&parmsz);

cl_device_id* cldevs = (cl_device_id *) malloc(parmsz);

clerr = clGetContextInfo(clctx, CL_CONTEXT_DEVICES, parmsz,
cldevs, NULL);

cl_command_queue clcmdq = clCreateCommandQueue(clctx, cldevs[0],
0, &clerr);
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## OpenCL Kernel Compilation: vadd

```
const char* vaddsrc =
    "__kernel void vadd(__global float *d_A, __global float *d_B,
__global float *d_C, int N) { \n"    [...etc and so forth...]
```

**OpenCL kernel source code as a big string**

```
cl_program clpgm;

clpgm = clCreateProgramWithSource(clctx, 1, &vaddsrc, NULL,
&clerr);
```

**Gives raw source code string(s) to OpenCL**

```
char clcompileflags[4096];
sprintf(clcompileflags, "-cl-mad-enable",
clerr = clBuildProgram(clpgm, 0, NULL, clcompileflags, NULL,
NULL);

cl_kernel clkern = clCreateKernel(clpgm, "vadd", &clerr);
```

**Set compiler flags, compile source, and retrieve a handle to the "vadd" kernel**

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## OpenCL Device Memory Allocation

### – **clCreateBuffer();**

- Allocates object in the device Global Memory
- Returns a pointer to the object
- Requires five parameters
  - OpenCL context pointer
  - Flags for access type by device (read/write, etc.)
  - Size of allocated object
  - Host memory pointer, if used in copy-from-host mode
  - Error code

### – **clReleaseMemObject()**

- Frees object
- Pointer to freed object

## OpenCL Device Memory Allocation (cont.)

### – **Code example:**

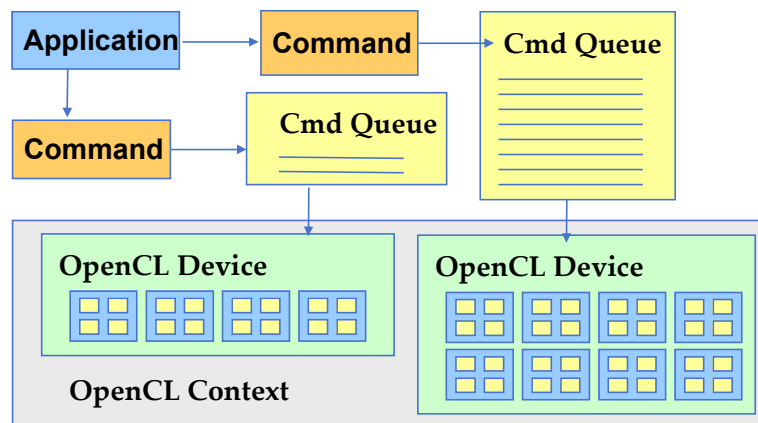
- Allocate a 1024 single precision float array
- Attach the allocated storage to d\_a
- “d\_” is often used to indicate a device data structure

```
VECTOR_SIZE = 1024;
cl_mem d_a;
int size = VECTOR_SIZE* sizeof(float);

d_a = clCreateBuffer(clctx, CL_MEM_READ_ONLY, size, NULL,
  NULL);
...
clReleaseMemObject(d_a);
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## OpenCL Device Command Execution



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## OpenCL Host-to-Device Data Transfer

- `clEnqueueWriteBuffer()` ;
  - Memory data transfer to device
  - Requires nine parameters
    - OpenCL command queue pointer
    - Destination OpenCL memory buffer
    - Blocking flag
    - Offset in bytes
    - Size (in bytes) of written data
    - Source host memory pointer
    - List of events to be completed before execution of this command
    - Event object tied to this command

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## OpenCL Device-to-Host Data Transfer

- `clEnqueueReadBuffer()` ;
  - Memory data transfer to host
  - requires nine parameters
    - OpenCL command queue pointer
    - Source OpenCL memory buffer
    - Blocking flag
    - Offset in bytes
    - Size of bytes of read data
    - Destination host memory pointer
    - List of events to be completed before execution of this command
    - Event object tied to this command

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## OpenCL Host-Device Data Transfer (cont.)

### –Code example:

- Transfer a  $64 * 64$  single precision float array
- `a` is in host memory and `d_a` is in device memory

```
clEnqueueWriteBuffer(clcmdq, d_a, CL_FALSE, 0,
                    mem_size, (const void *)a, 0, 0, NULL);
```

```
clEnqueueReadBuffer(clcmdq, d_result, CL_FALSE, 0,
                   mem_size, (void *) host_result, 0, 0, NULL);
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## OpenCL Host-Device Data Transfer (cont.)

- `clCreateBuffer` and `clEnqueueWriteBuffer` can be combined into a single command using special flags.
- Eg:
  - `d_A=clCreateBuffer(clctx,`
  - `CL_MEM_READ_ONLY | CL_MEM_COPY_HOST_PTR,`  
`mem_size, h_A, NULL);`
    - Combination of 2 flags here. `CL_MEM_COPY_HOST_PTR` to be used only if a valid host pointer is specified.
    - This creates a memory buffer on the device, and copies data from `h_A` into `d_A`.
    - Includes an implicit `clEnqueueWriteBuffer` operation, for all devices/command queues tied to the context `clctx`.

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Device Memory Allocation and Data Transfer for vadd

```
float *h_A = ..., *h_B = ...;
// allocate device (GPU) memory
cl_mem d_A, d_B, d_C;
d_A = clCreateBuffer(clctx, CL_MEM_READ_ONLY |
                     CL_MEM_COPY_HOST_PTR, N *sizeof(float), h_A,
NULL);
d_B = clCreateBuffer(clctx, CL_MEM_READ_ONLY |
                     CL_MEM_COPY_HOST_PTR, N *sizeof(float), h_B,
NULL);
d_C = clCreateBuffer(clctx, CL_MEM_WRITE_ONLY,
N *sizeof(float), NULL, NULL);
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Device Kernel Configuration Setting for vadd

```

clkern=clCreateKernel(clpgm, "vadd", NULL);

...

clerr= clSetKernelArg(clkern, 0,
sizeof(cl_mem), (void *)&d_A);

clerr= clSetKernelArg(clkern, 1,
sizeof(cl_mem), (void *)&d_B);

clerr= clSetKernelArg(clkern, 2,
sizeof(cl_mem), (void *)&d_C);

clerr= clSetKernelArg(clkern, 3, sizeof(int),
&N);

```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Device Kernel Launch and Remaining Code for vadd

```

cl_event event=NULL;
clerr= clEnqueueNDRangeKernel(clcmdq, clkern, 2, NULL,
    Gsz, Bsz, 0, NULL, &event);
clerr= clWaitForEvents(1, &event);
clEnqueueReadBuffer(clcmdq, d_C, CL_TRUE, 0,
    N*sizeof(float), h_C, 0, NULL, NULL);
clReleaseMemObject(d_A);
clReleaseMemObject(d_B);
clReleaseMemObject(d_C);
}

```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.



# 讲授内容

- **Related Programming Models: OpenCL**
- **Related Programming Models: OpenACC**
- **Multi-GPU: OpenMP**
- **Related Programming Models: MPI**

## 讲授内容： Related Programming Models: OpenACC

- ① **Introduction to OpenACC**
- ② **OpenACC Subtleties**

## OpenACC

- The OpenACC Application Programming Interface provides a set of
  - compiler directives (pragmas)
  - library routines and
  - environment variables

that can be used to write data parallel Fortran, C and C++ programs that run on accelerator devices including GPUs and CPUs

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## OpenACC Pragmas

- In C and C++, the **#pragma** directive is the method to provide to the compiler information that is not specified in the standard language.
  - These pragmas extend the base language

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Vector Addition in OpenACC

```
void VecAdd(float * __restrict__ output, const
           float * input1, const float * input 2, int
           inputLength)

{
    #pragma acc parallel loop
    copyin(input1[0:inputLength], input2[0:inputLength])
    copyout(output[0:inputLength])

    for(i = 0; i < inputLength; ++i) {
        output[i] = input1[i] + input2[i];
    }
}
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Simple Matrix-Matrix Multiplication in OpenACC

```
1. void computeAcc(float *P, const float *M, const float *N, int Mh, int Mw, int Nw)
2. {
3.     #pragma acc parallel loop copyin(M[0:Mh*Mw]) copyin(N[0:Mw*Nw]) copyout(P[0:Mh*Nw])
4.     for (int i=0; i<Mh; i++) {
5.         #pragma acc loop
6.         for (int j=0; j<Nw; j++) {
7.             float sum = 0;
8.             for (int k=0; k<Mw; k++) {
9.                 float a = M[i*Mw+k];
10.                float b = N[k*Nw+j];
11.                sum += a*b;
12.            }
13.            P[i*Nw+j] = sum;
14.        }
15.    }
16. }
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Some Observations (1)

```

1. void computeAcc(float *P, const float *M, const float *N, int Mh, int Mw, int Nw)
2. {
3.     #pragma acc parallel loop copyin(M[0:Mh*Mw]) copyin(N[0:Mw*Nw]) copyout(P[0:Mh*Nw])
4.     for (int i=0; i<Mh; i++) {
5.         #pragma acc loop
6.         for (int j=0; j<Nw; j++) {
7.             float sum = 0;
8.             for (int k=0; k<Mw; k++) {
9.                 float a = M[i*Mw+k];
10.                float b = N[k*Nw+j];
11.                sum += a*b;
12.            }
13.            P[i*Nw+j] = sum;
14.        }
15.    }
16. }

```

**The code is almost identical to the sequential version, except for the two lines with #pragma at line 3 and line 5.**

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Some Observations (2)

```

1. void computeAcc(float *P, const float *M, const float *N, int Mh, int Mw, int Nw)
2. {
3.     #pragma acc parallel loop copyin(M[0:Mh*Mw]) copyin(N[0:Mw*Nw]) copyout(P[0:Mh*Nw])
4.     for (int i=0; i<Mh; i++) {
5.         #pragma acc loop
6.         for (int j=0; j<Nw; j++) {
7.             float sum = 0;
8.             for (int k=0; k<Mw; k++) {
9.                 float a = M[i*Mw+k];
10.                float b = N[k*Nw+j];
11.                sum += a*b;
12.            }
13.            P[i*Nw+j] = sum;
14.        }
15.    }
16. }

```

**The #pragma at line 3 tells the compiler to generate code for the 'i' loop at line 4 through 15 so that the loop iterations are executed at the first level of parallelism on the accelerator.**

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Some Observations (3)

```

1. void computeAcc(float *P, const float *M, const float *N, int Mh, int Mw, int Nw)
2. {
3.     #pragma acc parallel loop copyin(M[0:Mh*Mw]) copyin(N[0:Mw*Nw]) copyout(P[0:Mh*Nw])
4.     for (int i=0; i<Mh; i++) {
5.         #pragma acc loop
6.         for (int j=0; j<Nw; j++) {
7.             float sum = 0;
8.             for (int k=0; k<Mw; k++) {
9.                 float a = M[i*Mw+k];
10.                float b = N[k*Nw+j];
11.                sum += a*b;
12.            }
13.            P[i*Nw+j] = sum;
14.        }
15.    }
16. }

```

The `copyin()` clause and the `copyout()` clause specify how the compiler should arrange for the matrix data to be transferred between the host and the accelerator.

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Some Observations (4)

```

1. void computeAcc(float *P, const float *M, const float *N, int Mh, int Mw, int Nw)
2. {
3.     #pragma acc parallel loop copyin(M[0:Mh*Mw]) copyin(N[0:Mw*Nw])
        copyout(P[0:Mh*Nw])
4.     for (int i=0; i<Mh; i++) {
5.         #pragma acc loop
6.         for (int j=0; j<Nw; j++) {
7.             float sum = 0;
8.             for (int k=0; k<Mw; k++) {
9.                 float a = M[i*Mw+k];
10.                float b = N[k*Nw+j];
11.                sum += a*b;
12.            }
13.            P[i*Nw+j] = sum;
14.        }
15.    }
16. }

```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Motivation

- **OpenACC programmers can often start with writing a sequential version and then annotate their sequential program with OpenACC directives.**
  - **leave most of the details in generating a kernel, memory allocation, and data transfers to the OpenACC compiler.**
- **OpenACC code can be compiled by non-OpenACC compilers by ignoring the pragmas.**


The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Frequently Encountered Issues


- **Some OpenACC pragmas are hints to the OpenACC compiler, which may or may not be able to act accordingly**
  - **The performance of an OpenACC program depends heavily on the quality of the compiler.**
  - **It may be hard to figure out why the compiler cannot act according to your hints**
  - **The uncertainty is much less so for CUDA or OpenCL programs**

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## OpenACC Device Model



## OpenACC Execution Model



# 讲授内容： Related Programming Models: OpenACC

## ① Introduction to OpenACC

## ② OpenACC Subtleties

### Parallel vs. Loop Constructs

```
#pragma acc parallel loop copyin(M[0:Mh*Mw]) copyin(N[0:Mw*Nw])
copyout(P[0:Mh*Nw])
```

```
for (int i=0; i<Mh; i++) {
```

```
...
```

```
}
```

is equivalent to:

```
#pragma acc parallel copyin(M[0:Mh*Mw]) copyin(N[0:Mw*Nw]) copyout(P[0:Mh*Nw])
```

```
{
```

```
    #pragma acc loop
```

```
    for (int i=0; i<Mh; i++) {
```

```
        ...
```

```
    }
```

```
}
```

(a parallel region that consists of a single loop)

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.



## More on Parallel Construct

```
#pragma acc parallel copyout(a) num_gangs(1024)
num_workers(32)
```

```
{
  a = 23;
}
```

**1024\*32 workers will be created. a=23 will be executed redundantly by all 1024 gang leads**

- A parallel construct is executed on an accelerator
- One can specify the number of gangs and number of workers in each gang
  - Equivalent to CUDA blocks and threads

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## What Does Each “Gang Loop” Do?

```
#pragma acc parallel
num_gangs(1024)
```

```
{
  for (int i=0; i<2048; i++) {
    ...
  }
}
```

```
#pragma acc parallel
num_gangs(1024)
```

```
{
  #pragma acc loop gang
  for (int i=0; i<2048; i++) {
    ...
  }
}
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Worker Loop

```
#pragma acc parallel num_gangs(1024) num_workers(32)
{
    #pragma acc loop gang
    for (int i=0; i<2048; i++) {
        #pragma acc loop worker
        for (int j=0; j<512; j++) {
            foo(i,j);
        }
    }
}
```

1024\*32=32K workers will be created, each executing 1M/32K = 32 instance of foo()

## A More Substantial Example

– Statements 1, 3, 5, 6 are redundantly executed by 32 gangs

<pre>#pragma acc parallel num_gangs(32) {     Statement 1;     #pragma acc loop gang     for (int i=0; i&lt;n; i++) {         Statement 2;     }     . . . }</pre>	<pre>#pragma acc parallel num_gangs(32) {     . . .     Statement 3;     #pragma acc loop gang     for (int i=0; i&lt;m; i++) {         Statement 4;     }     Statement 5;     if (condition) Statement 6; }</pre>
--	---

## A More Substantial Example

- The iterations of the  $n$  and  $m$  for-loop iterations are distributed to 32 gangs
- Each gang could further distribute the iterations to its workers
  - The number of workers in each gang will be determined by the compiler/runtime

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## A More Substantial Example

```
#pragma acc parallel
num_gangs(32)
{
    Statement 1;
    #pragma acc loop gang
    for (int i=0; i<n; i++) {
        Statement 2;
    }
    . . .
}
```

```
#pragma acc parallel num_gangs(32)
{
    . . .
    Statement 3;
    #pragma acc loop gang
    for (int i=0; i<m; i++) {
        Statement 4;
    }
    Statement 5;
    if (condition) Statement 6;
}
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Avoiding Redundant Execution

- Statements 1, 3, 5, 6 will be executed only once
- Iterations of the n and m loops will be distributed to 32 workers

```
#pragma acc parallel num_gangs(1)
num_workers(32)
{
    Statement 1;
    #pragma acc loop worker
    for (int i=0; i<n; i++) {
        Statement 2;
    }
    Statement 3;
    #pragma acc loop worker
    for (int i=0; i<m; i++) {
        Statement 4;
    }
    Statement 5;
    if (condition) Statement 6;
}
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Kernel Regions

- Kernel constructs are descriptive of programmer intentions
  - The compiler has a lot of flexibility in its use of the information
- This is in contrast with Parallel, which is prescriptive of the action for the compiler to follow

```
#pragma acc kernels
{
    #pragma acc loop gang(1024)
    for (int i=0; i<2048; i++) {
        a[i] = b[i];
    }
    #pragma acc loop gang(512)
    for (int j=0; j<2048; j++) {
        c[j] = a[j]*2;
    }
    for (int k=0; k<2048; k++) {
        d[k] = c[k];
    }
}
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Kernel Regions

- Code in a kernel region can be broken into multiple CUDA/OpenCL kernels
- The i, j, k loops can each become a kernel
  - The k-loop may even remain as host code
- Each kernel can have a different gang/worker configuration

```
#pragma acc kernels
{
    #pragma acc loop gang(1024)
    for (int i=0; i<2048; i++) {
        a[i] = b[i];
    }
    #pragma acc loop gang(512)
    for (int j=0; j<2048; j++) {
        c[j] = a[j]*2;
    }
    for (int k=0; k<2048; k++) {
        d[k] = c[k];
    }
}
```

## 讲授内容

- Related Programming Models: OpenCL
- Related Programming Models: OpenACC
- Multi-GPU: OpenMP
- Related Programming Models: MPI

# 讲授内容： Multi-GPU: OpenMP

## ① OpenMP

### ② Multi GPU Introduction I

### ③ Multi GPU Introduction II

### ④ Multi GPU patterns with OpenMP & Cooperative Groups

## OpenMP

- A parallel programming model based on the concepts of multithreading and shared memory
- OpenMP programs consists on annotations written into the serial code, called directives
- Needs a compiler with OpenMP support, which will translate the directives to the actual parallel code
- Highly portable across systems
- Limited to a single computer, however it can handle several processors

# Differences between process and threaded parallelism


## Process parallelism

- Uses local memory forcing each parallel worker to have its own memory space.
  - It needs explicit data sharing routines creating communication overhead.
- It's controlled by a main process.
- It can scale to multiple computers.
- Example: MPI (see Module 18)



## Threaded parallelism

- Uses shared memory, enabling all parallel workers to access the same memory space.
- It's controlled by a main thread, with all threads running on a single process.
- It can't scale beyond a single computer.
- Example: OpenMP



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

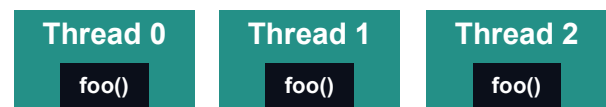
# OpenMP parallel regions

- OpenMP is a directive based parallel programming model, meaning we annotate the code to introduce parallelism.
- The directive to create a parallel region is:

- **#pragma omp parallel**
  - This pragma is added right before the block of code to parallelize.
  - This pragma runs the instructions inside the code block in each parallel thread.

–Example parallel behavior of calling a function inside a parallel region:

```
#pragma omp parallel
{
    foo();
}
```



Visualization of the behavior of the parallel region

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Parallelizing a “for” loop with OpenMP

- OpenMP is aware that in many cases “for” loops are a source for introducing parallelism, that is why it has special directives for loops
- To parallelize a for loop exists the directive:
  - **#pragma omp parallel for**
    - This pragma is added right before the for loop to parallelize.
    - An OpenMP compiler will identify the pragma and create a program where the loop is executed in parallel.

–Serial version of the vector sum algorithm:

```
for(int i = 0; i < n; ++i) {
    c[i] = a[i] + b[i];
}
```

–Parallel OpenMP version of the vector sum algorithm:

```
#pragma omp parallel for
for(int i = 0; i < n; ++i) {
    c[i] = a[i] + b[i];
}
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## #pragma omp parallel for

- Each variable used inside the loop has a kind determining if it's accessible by other threads: **shared**, **private**, **firstprivate**, **lastprivate**
- Shared variables are accessible by every thread in the region
- Private variables are local to each thread and are not initialized
- Firstprivate behave similarly to private variables but are initialized with the value before the parallel region
- Unless otherwise stated, every variable outside the loop will be considered shared and every inside the loop will be considered private

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.



## OpenMP Synchronization

- By default every parallel region will wait until the region is finished before executing other instructions. To change this behavior use the **nowait** option when declaring the parallel region
- **#pragma omp barrier** specifies a synchronization point equivalent to **\_\_syncthreads()** in CUDA, the pragma is not associated to a block of code

```
#pragma omp parallel for
for(int i = 0; i < n; ++i) {
    c[i] = a[i] + b[i];
    ...
    #pragma omp barrier
    c[i] += i >= 1 ? c[i - 1] : 0;
}
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## OpenMP Sections

- **Critical sections:**
  - Are executed by only one thread at a time
  - Are associated to a block of code
  - **#pragma omp critical**
- **Single sections:**
  - Are executed by only one thread of the team
  - Are associated to a block of code
  - **#pragma omp single**

```
int a = 0;
int b = 0;
#pragma omp parallel num_threads(4)
{
    #pragma omp critical
    a += 1 ;
    #pragma omp single
    b += 1;
}
```

- a final value is 4, because each thread executed the section one at a time
- b final value is 1, because only one thread executed the section

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## OpenMP API

- OpenMP in addition to the directives have an API, to control or obtain certain runtime parameters.
- As is the case with CUDA, OpenMP allows users to set the number of threads, get the number of threads and uniquely identify each thread.

### • `omp_set_num_threads( int )`:

- Sets the desired number of threads to be used by the OpenMP runtime on subsequent parallel regions.

### • `omp_get_num_threads()`:

- Returns the number of threads being used in the parallel region.

### • `omp_get_thread_num()`:

- Returns a numeric identifier for the calling thread, with different threads having different identifiers.

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Matrix Vector Multiplication OpenMP example

```
void tileMultiply(double* A, double* x, double y, int rows, int cols, int tileSize){
    int rid = omp_get_thread_num() * tileSize;           // Get the first row in the tile to compute
    for(int r = rid; r < min(rows, rid + tileSize); ++r ) {    // Iterate through the rows in the tile
        double sum = 0;
        for(int c = 0; c < cols; ++c)
            sum += A[r * cols + c] * x[c];                // Accumulate the dot product between r-
row and x
        y[r] = sum;                                         // Store the result into the
result
    }
}
...
#pragma omp parallel                                       // Create a parallel region
{
    int tnum = omp_get_num_threads();                     // Get the number of threads executing the region
    int tileSize = (rows + tnum - 1) / tnum;              // Calculate the number of rows in a tile
    tileMultiply(A, x, t, rows, cols, tileSize);          // Dispatch the function calculating the
multiplication
}
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Compiling an OpenMP program

- OpenMP compilers typically will ignore the OpenMP pragmas unless we specify the OpenMP flag.
- To use the API functions you need to include in the appropriate file:
  - `#include <omp.h>`
- Examples:
  - GCC & Clang: to compile an OpenMP annotated source file we use:
    - `gcc <file to compile> -fopenmp`
    - `clang <file to compile> -fopenmp`
    - Additionally if there are multiple OpenMP libraries, you may specify which version to use, for example `-lgomp` uses the GNU implementation.
  - NVCC: to compile an OpenMP annotated source file we use:
    - `nvcc <options> -Xcompiler -fopenmp`

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## 讲授内容： Multi-GPU: OpenMP


① OpenMP

② Multi GPU Introduction I

③ Multi GPU Introduction II

④ Multi GPU patterns with OpenMP & Cooperative Groups


## Schematic of a Multi GPU system



- The figure represents a system with 2 CPU's and 4 GPU's.
- GPU's are numbered from 0 to n-1, where n is the number of GPU's.
- The CUDA driver always starts with a default active device.
- There are two broad types of Multi GPU communication:
  - Through the PCIe bus
  - Through NVLINK

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.


## CUDA host API calls for Multi GPU's



- **cudaSetDevice()**
  - Set GPU device to use for device code execution on the active host thread.
  - Requires one parameter:
    - An int with the device id number
  - This function doesn't affect other host threads, meaning that setting the device on one thread will not set the device in other host threads. Also doesn't affect previous async calls.

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## CUDA host API calls for Multi GPU's



### • `cudaGetDevice()`

- Get GPU device being currently used by the active host thread.
- Requires one parameter:
  - An int pointer to store the device id

### • `cudaGetDeviceCount()`

- Get the number of CUDA-capable devices in the system.
- Requires one parameter:
  - An int pointer to store the device count

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## CUDA host API calls for Memory allocation with Multiple GPU's

To allocate or associate memory with a specific device using non-Managed CUDA-API calls, it's necessary to call `cudaSetDevice()` before doing the allocation call.

### • `cudaMalloc()`

- Allocates an object in the device global memory
- Two parameters
  - Address of a pointer to the allocated object
  - Size of allocated object in terms of bytes

### • `cudaHostAlloc()`

- Allocates pinned memory on the host
- Three parameters
  - Address of pointer to the allocated memory
  - Size of the allocated memory in bytes
  - Host Alloc flags

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## CUDA host API calls for Memory allocation with Multiple GPU's Unified Memory

- If the flag `cudaDevAttrConcurrentManagedAccess` is set in all devices, then it's not necessary to call `cudaSetDevice` before the `cudaMallocManaged` call.
- If the flag is not set but devices can access each others memory, then calling `cudaSetDevice` before the `cudaMallocManaged` call will establish the context for the managed memory on the active device.
  - With other devices accessing the data via PCIe at reduced bandwidth.

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## CUDA runtime calls affected by `cudaSetDevice`

- If `cudaSetDevice()` was called before a kernel launching call, the kernel will execute in the active device.
  - It's crucial that every non managed memory being used in the kernel resides in the active device, otherwise an error will occur.
- If `cudaSetDevice()` was called before a `cudaStreamCreate()`, then the stream will be associated with the active device.
- The synchronization functions: `cudaDeviceSynchronize()`, `cudaStreamSynchronize()` are also affected by `cudaSetDevice()`, synchronizing tasks only for the active device on the active host thread

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Putting it all together, vecAdd

```
float *m_A0, float *m_B0, *m_A1, float *m_B1, int n;
int size = n * sizeof(float);
cudaSetDevice(0); // Will set the active device to 0
cudaMalloc((void**) &m_A0, size); // Will allocate memory on device 0
cudaMalloc((void**) &m_B0, size); // Will allocate memory on device 0
cudaSetDevice(1); // Will set the active device to 1
cudaMalloc((void**) &m_A1, size); // Will allocate memory on device 1
cudaMalloc((void**) &m_B1, size); // Will allocate memory on device 1

// Memory initialization on the Host and memory transfers
cudaSetDevice(0); // Set the device for kernel
execution

vecAdd<<<gridDim, blockDim>>>(m_A0,m_B0);

cudaSetDevice(1); // Set the device for kernel
execution

vecAdd<<< gridDim, blockDim>>>(m_A1,m_B1);

cudaFree(m_A0); cudaFree(m_B0);
cudaFree(m_A1); cudaFree(m_B1);
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## 讲授内容： Multi-GPU: OpenMP


① OpenMP

② Multi GPU Introduction I

③ Multi GPU Introduction II

④ Multi GPU patterns with OpenMP & Cooperative Groups


## Memory transfers in a Multi GPU setup



- Involves transferring memory regions from one device to another, e.g. GPU0 to GPU1.
- There are three ways to do it:
  - Fully explicit memory transfers using `cudaMemcpyPeerAsync`, which requires the specification of the peer devices.
  - Partially explicit memory transfers using `cudaMemcpy`, relying on the unified address system.
  - Implicit peer memory access performed by the driver, without the need of explicit transfers.
- Not all three possibilities are available in every system.

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Explicit peer memory transfers CUDA host API functions



- `cudaMemcpyPeerAsync()`
    - Six parameters
      - Pointer to destination region on the destination device
      - Destination device id
      - Pointer to source region on the source device
      - Source device id
      - Number of bytes copied
      - CUDA stream
- Transfer between devices is asynchronous

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.



## Example: peer transfer cudaMemcpyPeerAsync

```
float *A0, *A1;

int size;

cudaSetDevice(0); // Set active device to 0
cudaMalloc((void**) &A0, size); // Allocate memory on device 0
cudaSetDevice(1); // Set active device to 1
cudaMalloc((void**) &A1, size); // Allocate memory on device 1


// Initialize region A0 on device 0

cudaMemcpyPeerAsync(A1, 1, A0, 0, size, stream); // Copy the data on A0 on device 0 to the region
A1 on device 1

cudaSetDevice(1); // Set the device for kernel execution
kernel<<<gridDim, blockDim, 0, stream>>>(A1); // Perform computations on A1

cudaFree(A0); // Free A0 region
cudaFree(A1); // Free A1 region
```

## Explicit peer memory transfers CUDA host API functions



- If the flag **cudaDevAttrUnifiedAddressing** is set to 1, then you may copy regions between devices using the traditional **cudaMemcpy** API function, setting the copy kind to **cudaMemcpyDefault**.
- To check if the flag is set you can use the API function:
  - **cudaDeviceGetAttribute()**

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Example: peer transfer cudaMemcpy

```
float *A0, *A1;

int size;

int unifiedAddr_flag0 = 0;

int unifiedAddr_flag1 = 0;
cudaSetDevice(0); // Set active device to 0
cudaMalloc((void**) &A0, size); // Allocate memory on device 0
cudaSetDevice(1); // Set active device to 1

cudaMalloc((void**) &A1, size); // Allocate memory on device 1

// Initialize region A0 on device 0

cudaDeviceGetAttribute(unifiedAddr_flag0, cudaDevAttrUnifiedAddressing, 0); // Check if unified
addressing is available on dev 0

cudaDeviceGetAttribute(unifiedAddr_flag1, cudaDevAttrUnifiedAddressing, 1); // Check if unified
addressing is available on dev 0


if( unified_addressing_flag0 == 1 && unified_addressing_flag1 == 1 )

    cudaMemcpy(A1, A0, size, cudaMemcpyDefault); // Copy the data on A0 on device 0 to the region A1 on
device 1

else

    // Throw error indicating the copy couldn't be performed
```

## Implicit peer memory access



- To query if implicit peer memory access are enabled use:
- `cudaDeviceCanAccessPeer`:
  - Three parameters:
    - Int pointer to place to store the flag.
    - Device id of device trying to access peer
    - Device id of peer device
  - This call is not symmetric, meaning that if the `canAccess` flag is set to 1 for deviceA and deviceB, it may not be set to 1 for deviceB and deviceA.


## Enabling and disabling implicit peer memory access

### • `cudaDeviceEnablePeerAccess`:

- Two parameters:
  - Device id to enable access to from the current active device.
  - Int flag set to 0.
- This call is not symmetric.
- Returns error `cudaErrorInvalidDevice` if not possible.

### • `cudaDeviceDisablePeerAccess`:

- One parameter:
  - Device id to disable access to from the current device.
- This call is not symmetric.



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Example: implicit peer access

```
float *ptrA; // Pointer to memory region on device devA
int devA;
int devB;
int BcanAccessA = 0;
cudaError_t error;


cudaDeviceCanAccessPeer(&BcanAccessA, devB, devA); // Check if devB can access devA memory.

cudaSetDevice(devB); // Set the current active device to devB
if(BcanAccessA == 0)
{
    error = cudaDeviceEnablePeerAccess(devA, 0); // Enable peer accesses to devA memory
    if(error == cudaSuccess) {
        kernel<<<gridDim, blockDim, 0, stream>>>(ptrA); // Access ptrA on device devA from device devB
    }
    cudaDeviceDisablePeerAccess(devA); // Disable peer access to devA, this call is not needed.
}
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## NVLINK

- Is a proprietary interconnect technology developed by NVIDIA
- Provides higher memory bandwidth communication between GPUS than PCIe communication
- NVLINK on the Tesla V100 delivers a 300 GB/s communication data rate, whereas the typical PCIe 3.0 link delivers only 32 GB/s



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.


## 讲授内容： Multi-GPU: OpenMP

- ① OpenMP
- ② Multi GPU Introduction I
- ③ Multi GPU Introduction II
- ④ Multi GPU patterns with OpenMP & Cooperative Groups

## Common parallel patterns in a Multi-GPU environment


### Batch processing:

- Execute the same independent task multiple times with different data.



### Cooperative patterns:


- Tasks need to cooperate between each other to collectively reach a goal.



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Batch processing

- Is an embarrassingly parallel pattern.
- With enough data it is usual we can achieve 100% usage of the compute resources.
- It's common on video and image processing applications, where we need to apply the same operation to lots of different data.




The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Batch processing

Typical operations to accomplish batch processing:

1. Get number of available devices.
  2. Considering the number of devices and number of desired tasks allocate, initialize and copy the memory need it by the algorithm.
  3. Create CUDA streams for each of the tasks to be executed concurrently.
  4. Launch in parallel the kernel.
- \* Remember to set the device at the beginning of each group of operations.



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Batch processing example with OpenMP


```

int deviceCount;
cudaGetDeviceCount(& deviceCount);

std::vector<cudaStream_t> streams(deviceCount);

#pragma omp parallel for num_threads(deviceCount)
for(int dev = 0; dev < deviceCount; ++dev) {
    cudaSetDevice(dev);
    cudaStreamCreate(&streams[i]);
    // Allocate, initialize and transfer memory
}


#pragma omp parallel for num_threads(deviceCount)
for(int dev = 0; dev < deviceCount; ++dev) {
    cudaSetDevice(dev);
    kernel<<<gridDim, blockDim,
streams[i]>>> (...);
}
  
```



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Cooperative patterns


- Might have unavoidable syncing points causing tasks to wait and thus wasting compute resources.
- In some cases even when massive amounts of input data it might not reach 100% resource usage.
- It's common on applications with steps to reach a goal like iterative algorithms.



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Cooperative patterns


- With cooperative patterns there is no single fit solution like with batch processing.
- Thus the process consists in a loop of:
  1. Launching the code in parallel.
  2. Profiling it.
  3. Analyzing and removing bottlenecks.



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Syncing patterns on different streams with OpenMP


```
std::vector<cudaStream_t> streams;
// Initialization of the streams on each
device.
#pragma omp parallel
{
    // Launch the different kernels on the
    streams.
    #pragma omp for num_threads(streams.size())
    for(auto& stream : streams)
        cudaStreamSynchronize(stream);
    #pragma omp barrier
}
```



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Multi-GPU Syncing patterns with Cooperative Groups

- Cooperative Groups is a C++-CUDA high level abstraction to perform syncing across different parallel granularities (Threads, Blocks, Grids, and Devices).
- Multi-GPU syncing with cooperatives groups requires:
  - Devices with the exact same compute capability.
  - Compute capability of 6 or higher.
  - Executing the same kernel across all devices.




The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.



## Multi-GPU Syncing patterns with Cooperative Groups

- To enable the use of Cooperative Groups we need to include the file **cooperative\_groups.h** and use the namespace **cooperative\_groups**.
- The kernels need to be compiled using separate compilation and then linked with the **-rdc=true** flag.
- You also need to ensure that MPS is disabled and **CU\_DEVICE\_ATTRIBUTE\_COOPERATIVE\_MULTI\_DEVICE\_LAUNCH** is set in the device properties using **cuDeviceGetAttribute** the API function.



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Multi-GPU Syncing patterns with Cooperative Groups

Launching a kernel with Multi-GPU syncing and Cooperative Groups requires using the API function:

- **cudaLaunchCooperativeKernelMultiDevice**:
  - The first parameter is an array of **CUDA\_LAUNCH\_PARAMS**, where except for kernel params and the stream, all other fields need to be the same.
  - The number of devices to use.

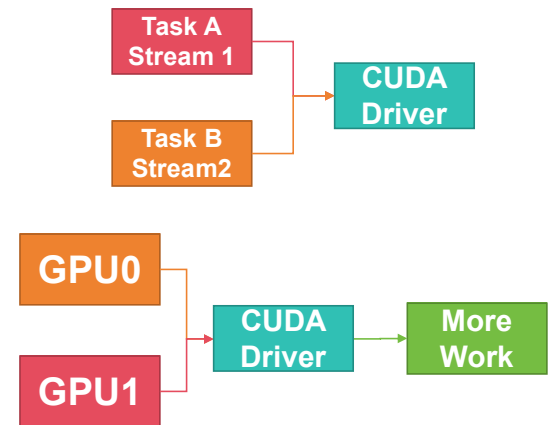
```

typedef struct CUDA_LAUNCH_PARAMS_st {
    CUfunction function;           // Kernel to launch.
    unsigned int gridDimX;        // Grid dimensions.
    unsigned int gridDimY;
    unsigned int gridDimZ;
    unsigned int blockDimX;       // Block dimensions.
    unsigned int blockDimY;
    unsigned int blockDimZ;
    unsigned int sharedMemBytes;  // Shared memory size.
    CUstream hStream;             // Stream to perform the work
    void **kernelParams;          // Kernel parameters
} CUDA_LAUNCH_PARAMS;
  
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Multi-GPU Syncing patterns with Cooperative Groups

```
#include <cooperative_groups.h>
using namespace cooperative_groups;
void __global__ kernel(...) {
    // Work
    multi_grid_group multi_grid =
this_multi_grid();
    multi_grid.sync();
    // Work
}
// Work
cudaLaunchCooperativeKernelMultiDevice
(...);
```



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## 讲授内容

- Related Programming Models: OpenCL
- Related Programming Models: OpenACC
- Multi-GPU: OpenMP
- Related Programming Models: MPI

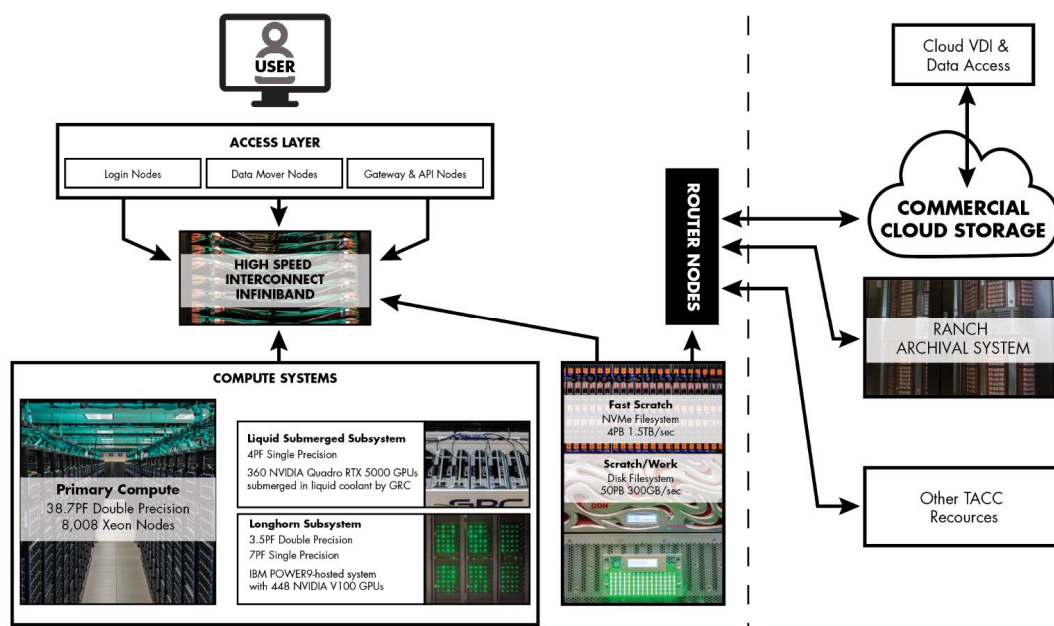
# 讲授内容： Related Programming Models: MPI

## ① Warps and SIMD Hardware

## ② Performance Impact of Control Divergence


## ③ Overlapping Computation with Communication

### Frontera – Installed at TACC 9/2019



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## MPI – Programming and Execution Model



**Many processes distributed in a cluster**  
**Each process computes part of the output**  
**Processes communicate with each other**  
**Processes can synchronize**

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## MPI Initialization, Info and Sync

```
int MPI_Init(int *argc, char ***argv)
```

Initialize MPI

```
MPI_COMM_WORLD
```

MPI group with all allocated nodes

```
int MPI_Comm_rank (MPI_Comm comm, int *rank)
```

Rank of the calling process in group of comm

```
int MPI_Comm_size (MPI_Comm comm, int *size)
```

Number of processes in the group of comm

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Vector Addition: Main Process

```
int main(int argc, char *argv[]) {
    int vector_size = 1024 * 1024 * 1024;
    int pid=-1, np=-1;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &pid);
    MPI_Comm_size(MPI_COMM_WORLD, &np);

    if(np < 3) {
        if(0 == pid) printf("Need 3 or more processes.\n");
        MPI_Abort( MPI_COMM_WORLD, 1 ); return 1;
    }
```



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Vector Addition: Main Process

```
if(pid < np - 1)
    compute_node(vector_size / (np - 1));
else
    data_server(vector_size);

MPI_Finalize();
return 0;
}
```



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## MPI Sending Data

```
int MPI_Send(void *buf, int count, MPI_Datatype
datatype, int dest, int tag, MPI_Comm comm)
```

**Buf:** Initial address of send buffer (choice)

**Count:** Number of elements in send buffer (nonnegative integer)

**Datatype:** Datatype of each send buffer element (handle)


**Dest:** Rank of destination (integer)

**Tag:** Message tag (integer)

**Comm:** Communicator (handle)

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## MPI Send



```
int MPI_Send(void *buf, int count, MPI_Datatype
datatype, int dest, int tag, MPI_Comm comm)
```

**Buf:** Initial address of send buffer (choice)

**Count:** Number of elements in send buffer (nonnegative integer)

**Datatype:** Datatype of each send buffer element (handle)

**Dest:** Rank of destination (integer)

**Tag:** Message tag (integer)

**Comm:** Communicator (handle)

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## MPI Receiving Data

```
int MPI_Recv(void *buf, int count, MPI_Datatype
datatype, int source, int tag, MPI_Comm comm,
MPI_Status *status)
```

**Buf:** Initial address of receive buffer (choice)

**Count:** Maximum number of elements in receive buffer (integer)

**Datatype:** Datatype of each receive buffer element (handle)

**Source:** Rank of source (integer)


**Tag:** Message tag (integer)

**Comm:** Communicator (handle)

**Status:** Status object (Status)

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## MPI Recv



```
int MPI_Recv(void *buf, int count, MPI_Datatype datatype,
int source, int tag, MPI_Comm comm, MPI_Status *status)
```

**Buf:** Initial address of receive buffer (choice)

**Count:** Maximum number of elements in receive buffer (integer)

**Datatype:** Datatype of each receive buffer element (handle)

**Source:** Rank of source (integer)

**Tag:** Message tag (integer)

**Comm:** Communicator (handle)

**Status:** Status object (Status)

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Vector Addition: Server Process (I)

```
void data_server(unsigned int vector_size) {
    int np, num_nodes = np - 1, first_node = 0, last_node = np - 2;
    unsigned int num_bytes = vector_size * sizeof(float);
    float *input_a = 0, *input_b = 0, *output = 0;

    /* Set MPI Communication Size */
    MPI_Comm_size(MPI_COMM_WORLD, &np);

    /* Allocate input data */
    input_a = (float *)malloc(num_bytes);
    input_b = (float *)malloc(num_bytes);
    output = (float *)malloc(num_bytes);
    if(input_a == NULL || input_b == NULL || output == NULL) {
        printf("Server couldn't allocate memory\n");
        MPI_Abort( MPI_COMM_WORLD, 1 );
    }
    /* Initialize input data */
    random_data(input_a, vector_size, 1, 10);
    random_data(input_b, vector_size, 1, 10);
}
```

## Vector Addition: Server Process (II)

```
/* Send data to compute nodes */
float *ptr_a = input_a;
float *ptr_b = input_b;
for(int process = 1; process < last_node; process++) {
    MPI_Send(ptr_a, vector_size / num_nodes, MPI_FLOAT,
             process, DATA_DISTRIBUTE, MPI_COMM_WORLD);
    ptr_a += vector_size / num_nodes;

    MPI_Send(ptr_b, vector_size / num_nodes, MPI_FLOAT,
             process, DATA_DISTRIBUTE, MPI_COMM_WORLD);
    ptr_b += vector_size / num_nodes;
}
/* Wait for nodes to compute */
MPI_Barrier(MPI_COMM_WORLD);
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.



## Vector Addition: Server Process (III)

```

/* Wait for previous communications */
MPI_Barrier(MPI_COMM_WORLD);
/* Collect output data */
MPI_Status status;
for(int process = 0; process < num_nodes;
process++) {
    MPI_Recv(output + process * num_points /
num_nodes,
            num_points / num_comp_nodes, MPI_REAL, process,
            DATA_COLLECT, MPI_COMM_WORLD, &status );
}
/* Store output data */
store_output(output, dimx, dimy, dimz);
/* Release resources */
free(input_a);
free(input_b);
free(output);
}

```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Vector Addition: Compute Process (I)

```

void compute_node(unsigned int vector_size ) {
    int np;
    unsigned int num_bytes = vector_size * sizeof(float);
    float *input_a, *input_b, *output;
    MPI_Status status;

    MPI_Comm_size(MPI_COMM_WORLD, &np);
    int server_process = np - 1;
    /* Alloc host memory */
    input_a = (float *)malloc(num_bytes);
    input_b = (float *)malloc(num_bytes);
    output = (float *)malloc(num_bytes);
    /* Get the input data from server process */
    MPI_Recv(input_a, vector_size, MPI_FLOAT, server_process,
            DATA_DISTRIBUTE, MPI_COMM_WORLD, &status);
    MPI_Recv(input_b, vector_size, MPI_FLOAT, server_process,
            DATA_DISTRIBUTE, MPI_COMM_WORLD, &status);
}

```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## MPI Barriers

- `int MPI_Barrier (MPI_Comm comm)`
  - **Comm: Communicator (handle)**
- **Blocks the caller until all group members have called it; the call returns at any process only after all group members have entered the call.**

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## MPI Barriers

- **Wait until all other processes in the MPI group reach the same barrier**


Example Code

```
Do_stuff();
```

```
MPI_Barrier();
```

```
Do_more_stuff();
```

- All processes are executing `Do_Stuff()`
- Some processes reach the barrier and the wait in the barrier until all reach the barrier



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Vector Addition: Compute Process (II)

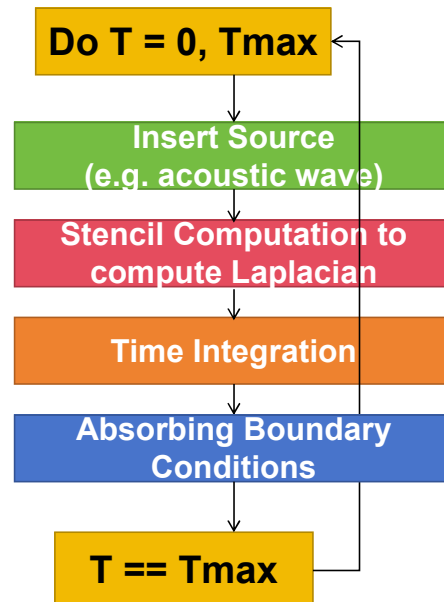
```
/* Compute the partial vector addition */
for(int i = 0; i < vector_size; ++i) {
    output[i] = input_a[i] + input_b[i];
}
/* Report to barrier after computation is done*/
MPI_Barrier(MPI_COMM_WORLD);
/* Send the output */
MPI_Send(output, vector_size, MPI_FLOAT,
         server_process, DATA_COLLECT, MPI_COMM_WORLD);
/* Release memory */
free(input_a);
free(input_b);
free(output);
}
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## 讲授内容：Related Programming Models: MPI

- ① Warps and SIMD Hardware
- ② Performance Impact of Control Divergence
- ③ Overlapping Computation with Communication


# A Typical Wave Propagation Application




The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Review of Stencil Computations

- Example: wave propagation modeling
- $\nabla^2 U - \frac{1}{v^2} \frac{\partial U}{\partial t} = 0$
- Approximate Laplacian using finite differences



Laplacian  
and Time  
Integration



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Wave Propagation: Kernel Code

```

/* Coefficients used to calculate the laplacian */
__constant__ float coeff[5];

__global__ void wave_propagation(float *next, float *in,
                                float *prev, float *velocity, dim3 dim)
{
    unsigned x = threadIdx.x + blockIdx.x * blockDim.x;
    unsigned y = threadIdx.y + blockIdx.y * blockDim.y;
    unsigned z = threadIdx.z + blockIdx.z * blockDim.z;

    /* Point index in the input and output matrixes */
    unsigned n = x + y * dim.z + z * dim.x * dim.y;

    /* Only compute for points within the matrixes */
    if(x < dim.x && y < dim.y && z < dim.z) {

        /* Calculate the contribution of each point to the laplacian */
        float laplacian = coeff[0] + in[n];

```

## Wave Propagation: Kernel Code

```

        for(int i = 1; i < 5; ++i) {
            laplacian += coeff[i] *
                (in[n - i] + /* Left */
                 in[n + i] + /* Right */
                 in[n - i * dim.x] + /* Top */
                 in[n + i * dim.x] + /* Bottom
*/
                 in[n - i * dim.x * dim.y] + /*
Behind */
                 in[n + i * dim.x * dim.y]); /*
Front */
        }

        /* Time integration */
        next[n] = velocity[n] * laplacian + 2 * in[n]
        - prev[n];
    }
}


```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Stencil Domain Decomposition

– Volumes are split into tiles (along the Z-axis)

– 3D-Stencil introduces data dependencies



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Wave Propagation: Main Process

```
int main(int argc, char *argv[]) {
    int pad = 0, dimx = 480+pad, dimy = 480, dimz = 400, nreps = 100;
    int pid=-1, np=-1;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &pid);
    MPI_Comm_size(MPI_COMM_WORLD, &np);

    if(np < 3) {
        if(0 == pid) printf("Needed 3 or more processes.\n");
        MPI_Abort( MPI_COMM_WORLD, 1 ); return 1;
    }
    if(pid < np - 1)
        compute_node(dimx, dimy, dimz / (np - 1), nreps);
    else
        data_server( dimx, dimy, dimz, nreps );

    MPI_Finalize();
    return 0;
}
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Stencil Code: Server Process (I)

```
void data_server(int dimx, int dimy, int dimz, int nreps)
{
    int np, num_comp_nodes = np - 1, first_node = 0,
    last_node = np - 2;
    unsigned int num_points = dimx * dimy * dimz;
    unsigned int num_bytes = num_points * sizeof(float);
    float *input=0, *output = NULL, *velocity = NULL;
    /* Set MPI Communication Size */
    MPI_Comm_size(MPI_COMM_WORLD, &np);
    /* Allocate input data */
    input = (float *)malloc(num_bytes);
    output = (float *)malloc(num_bytes);
    velocity = (float *)malloc(num_bytes);
    if(input == NULL || output == NULL || velocity == NULL)
    {
        printf("Server couldn't allocate memory\n");
        MPI_Abort( MPI_COMM_WORLD, 1 );
    }
    /* Initialize input data and velocity */
    random_data(input, dimx, dimy, dimz, 1, 10);
    random_data(velocity, dimx, dimy, dimz, 1, 10);
}
```

## Stencil Code: Server Process (II)

```
/* Calculate number of shared points */
int edge_num_points = dimx * dimy * (dimz / num_comp_nodes + 4);
int int_num_points = dimx * dimy * (dimz / num_comp_nodes + 8);
float *input_send_address = input;

/* Send input data to the first compute node */
MPI_Send(send_address, edge_num_points, MPI_REAL, first_node,
        DATA_DISTRIBUTE, MPI_COMM_WORLD );
send_address += dimx * dimy * (dimz / num_comp_nodes - 4);

/* Send input data to "internal" compute nodes */
for(int process = 1; process < last_node; process++) {
    MPI_Send(send_address, int_num_points, MPI_FLOAT, process,
            DATA_DISTRIBUTE, MPI_COMM_WORLD);
    send_address += dimx * dimy * (dimz / num_comp_nodes);
}

/* Send input data to the last compute node */
MPI_Send(send_address, edge_num_points, MPI_REAL, last_node,
        DATA_DISTRIBUTE, MPI_COMM_WORLD);
```

## Stencil Code: Server Process (II)

```
float *velocity_send_address = velocity;

/* Send velocity data to compute nodes */
for(int process = 0; process < last_node + 1; process++) {
    MPI_Send(send_address, edge_num_points, MPI_FLOAT, process,
             DATA_DISTRIBUTE, MPI_COMM_WORLD);
    send_address += dimx * dimy * (dimz / num_comp_nodes);
}

/* Wait for nodes to compute */
MPI_Barrier(MPI_COMM_WORLD);

/* Collect output data */
MPI_Status status;
for(int process = 0; process < num_comp_nodes; process++)
    MPI_Recv(output + process * num_points / num_comp_nodes,
             num_points / num_comp_nodes, MPI_FLOAT, process,
             DATA_COLLECT, MPI_COMM_WORLD, &status );
}
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Stencil Code: Server Process (III)

```
/* Store output data */
store_output(output, dimx,
dimy, dimz);

/* Release resources */
free(input);
free(velocity);
free(output);
}
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.



## Stencil Code: Compute Process (I)

```
void compute_node_stencil(int dimx, int dimy, int dimz, int nreps ) {
    int np, pid;
    MPI_Comm_rank(MPI_COMM_WORLD, &pid);
    MPI_Comm_size(MPI_COMM_WORLD, &np);

    unsigned int num_points      = dimx * dimy * (dimz + 8);
    unsigned int num_bytes      = num_points * sizeof(float);
    unsigned int num_ghost_points = 4 * dimx * dimy;
    unsigned int num_ghost_bytes = num_ghost_points * sizeof(float);

    int left_ghost_offset  = 0;
    int right_ghost_offset = dimx * dimy * (4 + dimz);

    float *input = NULL, *output = NULL, *prev = NULL, *v = NULL;

    /* Allocate device memory for input and output data */
    gmacMalloc((void **)&input,  num_bytes);
    gmacMalloc((void **)&output,  num_bytes);
    gmacMalloc((void **)&prev,   num_bytes);
    gmacMalloc((void **)&v,      num_bytes);
}
```

## Stencil Code: Compute Process (II)

```
MPI_Status status;
int left_neighbor  = (pid > 0) ? (pid - 1) : MPI_PROC_NULL;
int right_neighbor = (pid < np - 2) ? (pid + 1) : MPI_PROC_NULL;
int server_process = np - 1;

/* Get the input data from server process */
float *rcv_address = input + num_ghost_points * (0 == pid);
MPI_Recv(rcv_address, num_points, MPI_FLOAT, server_process,
        DATA_DISTRIBUTE, MPI_COMM_WORLD, &status );


/* Get the velocity data from server process */
rcv_address = h_v + num_ghost_points * (0 == pid);
MPI_Recv(rcv_address, num_points, MPI_FLOAT, server_process,
        DATA_DISTRIBUTE, MPI_COMM_WORLD, &status );
```

# 讲授内容： Related Programming Models: MPI

- ① Warps and SIMD Hardware
- ② Performance Impact of Control Divergence
- ③ Overlapping Computation with Communication

## Stencil Domain Decomposition

- Volumes are split into tiles (along the Z-axis)
- 3D-Stencil introduces data dependencies



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.


## CUDA and MPI Communication

### – Source MPI process:

- `cudaMemcpy(tmp,src, cudaMemcpyDeviceToHost)`
- `MPI_Send()`

### – Destination MPI process:

- `MPI_Recv()`
- `cudaMemcpy(dst, src, cudaMemcpyDeviceToDevice)`



## Data Server Process Code (I)

```
void data_server(int dimx, int dimy, int dimz, int nreps) {
    int np,
    /* Set MPI Communication Size */
    MPI_Comm_size(MPI_COMM_WORLD, &np);

    num_comp_nodes = np - 1, first_node = 0, last_node = np - 2;
    unsigned int num_points = dimx * dimy * dimz;
    unsigned int num_bytes = num_points * sizeof(float);
    float *input=0, *output=0;
    /* Allocate input data */
    input = (float *)malloc(num_bytes);
    output = (float *)malloc(num_bytes);
    if(input == NULL || output == NULL) {
        printf("server couldn't allocate memory\n");
        MPI_Abort( MPI_COMM_WORLD, 1 );
    }
    /* Initialize input data */
    random_data(input, dimx, dimy, dimz, 1, 10);
    /* Calculate number of shared points */
    int edge_num_points = dimx * dimy * (dimz / num_comp_nodes + 4);
    int int_num_points = dimx * dimy * (dimz / num_comp_nodes + 8);
    float *send_address = input;
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Data Server Process Code (II)

```

/* Send data to the first compute node */
MPI_Send(send_address, edge_num_points, MPI_FLOAT, first_node,
         0, MPI_COMM_WORLD );

send_address += dimx * dimy * (dimz / num_comp_nodes - 4);
/* Send data to "internal" compute nodes */
for(int process = 1; process < last_node; process++) {
    MPI_Send(send_address, int_num_points, MPI_FLOAT, process,
             0, MPI_COMM_WORLD);
    send_address += dimx * dimy * (dimz / num_comp_nodes);
}

/* Send data to the last compute node */
MPI_Send(send_address, edge_num_points, MPI_FLOAT, last_node,
         0, MPI_COMM_WORLD);

```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Compute Process Code (I).

```

void compute_node_stencil(int dimx, int dimy, int dimz, int nreps ) {
    int np, pid;
    MPI_Comm_rank(MPI_COMM_WORLD, &pid);
    MPI_Comm_size(MPI_COMM_WORLD, &np);
    int server_process = np - 1;

    unsigned int num_points      = dimx * dimy * (dimz + 8);
    unsigned int num_bytes      = num_points * sizeof(float);
    unsigned int num_halo_points = 4 * dimx * dimy;
    unsigned int num_halo_bytes = num_halo_points * sizeof(float);

    /* Alloc host memory */
    float *h_input = (float *)malloc(num_bytes);
    /* Alloc device memory for input and output data */
    float *d_input = NULL;
    cudaMalloc((void **)&d_input, num_bytes );
    float *rcv_address = h_input + num_halo_points * (0 == pid);
    MPI_Recv(rcv_address, num_points, MPI_FLOAT, server_process,
             MPI_ANY_TAG, MPI_COMM_WORLD, &status );
    cudaMemcpy(d_input, h_input, num_bytes, cudaMemcpyHostToDevice);
}

```

## Stencil Code: Kernel Launch

```
void launch_kernel(float *next, float *in, float
*prev, float *velocity,
                  int dimx, int dimy, int dimz)
{
    dim3 Gd, Bd, Vd;

    Vd.x = dimx; Vd.y = dimy; Vd.z = dimz;

    Bd.x = BLOCK_DIM_X; Bd.y = BLOCK_DIM_Y; Bd.z =
BLOCK_DIM_Z;

    Gd.x = (dimx + Bd.x - 1) / Bd.x;
    Gd.y = (dimy + Bd.y - 1) / Bd.y;
    Gd.z = (dimz + Bd.z - 1) / Bd.z;

    wave_propagation<<<Gd, Bd>>>(next, in, prev,
velocity, Vd);
}
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## MPI Sending and Receiving Data

- **int MPI\_Sendrecv(void \*sendbuf, int sendcount, MPI\_Datatype sendtype, int dest, int sendtag, void \*recvbuf, int recvcount, MPI\_Datatype recvtype, int source, int recvtag, MPI\_Comm comm, MPI\_Status \*status)**
  - **Sendbuf:** Initial address of send buffer (choice)
  - **Sendcount:** Number of elements in send buffer (integer)
  - **Sendtype:** Type of elements in send buffer (handle)
  - **Dest:** Rank of destination (integer)
  - **Sendtag:** Send tag (integer)
  - **Recvcount:** Number of elements in receive buffer (integer)
  - **Recvtype:** Type of elements in receive buffer (handle)
  - **Source:** Rank of source (integer)
  - **Recvtag:** Receive tag (integer)
  - **Comm:** Communicator (handle)
  - **Recvbuf:** Initial address of receive buffer (choice)
  - **Status:** Status object (Status). This refers to the receive operation.

## Compute Process Code (II)

```
float *h_output = NULL, *d_output = NULL, *d_vsqr = NULL;
float *h_output = (float *)malloc(num_bytes);
cudaMalloc((void **)&d_output, num_bytes );
```

```
float *h_left_boundary = NULL, *h_right_boundary = NULL;
float *h_left_halo = NULL, *h_right_halo = NULL;
```

```
/* Alloc host memory for halo data */
```


```
cudaHostAlloc((void **)&h_left_boundary, num_halo_bytes, cudaHostAllocDefault);
cudaHostAlloc((void **)&h_right_boundary, num_halo_bytes, cudaHostAllocDefault);
cudaHostAlloc((void **)&h_left_halo, num_halo_bytes, cudaHostAllocDefault);
cudaHostAlloc((void **)&h_right_halo, num_halo_bytes, cudaHostAllocDefault);
```

```
/* Create streams used for stencil computation */
```

```
cudaStream_t stream0, stream1;
cudaStreamCreate(&stream0);
cudaStreamCreate(&stream1);
```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Device Memory Offsets Used for Data Exchange with Neighbors



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Compute Process Code (III)

```

MPI_Status status;
int left_neighbor = (pid > 0) ? (pid - 1) : MPI_PROC_NULL;
int right_neighbor = (pid < np - 2) ? (pid + 1) : MPI_PROC_NULL;

/* Upload stencil coefficients */
upload_coefficients(coeff, 5);

int left_halo_offset = 0;
int right_halo_offset = dimx * dimy * (4 + dimz);
int left_stagel_offset = 0;
int right_stagel_offset = dimx * dimy * (dimz - 4);
int stage2_offset = num_halo_points;

MPI_Barrier( MPI_COMM_WORLD );
for(int i=0; i < nreps; i++) {
    /* Compute boundary values needed by other nodes first */
    launch_kernel(d_output + left_stagel_offset,
                  d_input + left_stagel_offset, dimx, dimy, 12, stream0);
    launch_kernel(d_output + right_stagel_offset,
                  d_input + right_stagel_offset, dimx, dimy, 12, stream0);

    /* Compute the remaining points */
    launch_kernel(d_output + stage2_offset, d_input + stage2_offset,
                  dimx, dimy, dimz, stream1);
}

```

## Compute Process Code (IV)

```

/* Copy the data needed by other nodes to the
host */
    cudaMemcpyAsync(h_left_boundary, d_output +
num_halo_points,
                    num_halo_bytes, cudaMemcpyDeviceToHost,
stream0 );
    cudaMemcpyAsync(h_right_boundary,
                    d_output + right_stagel_offset +
num_halo_points,
                    num_halo_bytes, cudaMemcpyDeviceToHost,
stream0 );
    cudaStreamSynchronize(stream0);

```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Syntax for MPI\_Sendrecv()

- **int MPI\_Sendrecv(void \*sendbuf, int sendcount, MPI\_Datatype sendtype, int dest, int sendtag, void \*recvbuf, int recvcount, MPI\_Datatype recvtype, int source, int recvtag, MPI\_Comm comm, MPI\_Status \*status)**
  - **Sendbuf:** Initial address of send buffer (choice)
  - **Sendcount:** Number of elements in send buffer (integer)
  - **Sendtype:** Type of elements in send buffer (handle)
  - **Dest:** Rank of destination (integer)
  - **Sendtag:** Send tag (integer)
  - **Recvcount:** Number of elements in receive buffer (integer)
  - **Recvtype:** Type of elements in receive buffer (handle)
  - **Source:** Rank of source (integer)
  - **Recvtag:** Receive tag (integer)
  - **Comm:** Communicator (handle)
  - **Recvbuf:** Initial address of receive buffer (choice)
  - **Status:** Status object (Status). This refers to the receive operation.

## Compute Process Code (V)

```

/* Send data to left, get data from right */
MPI_Sendrecv(h_left_boundary, num_halo_points, MPI_FLOAT,
             left_neighbor, i, h_right_halo,
             num_halo_points, MPI_FLOAT, right_neighbor, i,
             MPI_COMM_WORLD, &status );
/* Send data to right, get data from left */
MPI_Sendrecv(h_right_boundary, num_halo_points, MPI_FLOAT,
             right_neighbor, i, h_left_halo,
             num_halo_points, MPI_FLOAT, left_neighbor, i,
             MPI_COMM_WORLD, &status );

cudaMemcpyAsync(d_output+left_halo_offset, h_left_halo,
               num_halo_bytes, cudaMemcpyHostToDevice, stream0);
cudaMemcpyAsync(d_output+right_ghost_offset, h_right_ghost,
               num_halo_bytes, cudaMemcpyHostToDevice, stream0 );
cudaDeviceSynchronize();
float *temp = d_output;
d_output = d_input; d_input = temp;
}

```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.



## Compute Process Code (VI)

```

/* Wait for previous communications */
MPI_Barrier(MPI_COMM_WORLD);

float *temp = d_output;
d_output = d_input;
d_input = temp;

/* Send the output, skipping halo points */
cudaMemcpy(h_output, d_output, num_bytes,
            cudaMemcpyDeviceToHost);
float *send_address = h_output + num_ghost_points;
MPI_Send(send_address, dimx * dimy * dimz, MPI_REAL,
          server_process, DATA_COLLECT, MPI_COMM_WORLD);
MPI_Barrier(MPI_COMM_WORLD);

/* Release resources */
free(h_input); free(h_output);
cudaFreeHost(h_left_ghost_own); cudaFreeHost(h_right_ghost_own);
cudaFreeHost(h_left_ghost); cudaFreeHost(h_right_ghost);
cudaFree(d_input); cudaFree(d_output);
}

```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## Data Server Code (III)

```

/* Wait for nodes to compute */
MPI_Barrier(MPI_COMM_WORLD);

/* Collect output data */
MPI_Status status;
for(int process = 0; process < num_comp_nodes; process++)
    MPI_Recv(output + process * num_points / num_comp_nodes,
             num_points / num_comp_nodes, MPI_REAL, process,
             DATA_COLLECT, MPI_COMM_WORLD, &status);

/* Store output data */
store_output(output, dimx, dimy, dimz);

/* Release resources */
free(input);
free(output);
}

```

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

## More on MPI Message Types

### —Point-to-point communication

- Send and Receive

### —Collective communication

- Barrier

- Broadcast

- Reduce

- Gather and Scatter

The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the Creative Commons Attribution-NonCommercial 4.0 International License.

**T H A N K S**